

Subverting System Authentication With Context-Aware, Reactive Virtual Machine Introspection

Yangchun Fu, Zhiqiang Lin, Kevin Hamlen

Department of Computer Science
The University of Texas at Dallas

December 12th, 2013

Malicious Virtual Machine Monitor (VMM)

Goal

- Subverting authentication (e.g., `login`) with Context-Aware, Reactive Virtual Machine Introspection(VMI)
- Attackers can gain both fun and profit, by accessing sensitive data in a computer

Malicious Virtual Machine Monitor (VMM)


Goal

- Subverting authentication (e.g., `login`) with Context-Aware, Reactive Virtual Machine Introspection(VMI)
- Attackers can gain both fun and profit, by accessing sensitive data in a computer


Assumptions

- Assume physical access (lost of laptop, VMs running in a cloud)
- Possible attackers/users
 - Malicious cloud providers (cloud being compromised)
 - Law enforcement (accessing criminal's computer)

Running a machine inside a malicious VMM




Running a machine inside a malicious VMM



Inception Attack

- Changing your idea using a dream
- Dream can be inside a dream

Running a machine inside a malicious VMM




Inception Attack

- Changing your idea using a dream
- Dream can be inside a dream


Malicious Virtualization Monitor

- Running a machine inside a **virtual machine**
- We change the guest OS state from the **malicious virtual machine** without the awareness from any inside programs

Traditional computer system structure




Traditional computer system structure



Authentication protection Mechanism

- Anti-debugging Logic
- Cryptographic Security
- Code Obfuscation
- Self-Checking

Traditional computer system structure




Trust?


Authentication protection Mechanism

- anti-debugging logic
- cryptographic security
- code obfuscation
- self-checking

Virtualization



Key Idea



Adding a virtualization layer

- VMM runs at higher privilege than guest OS
- Great isolation, more stealthy
- A full control of guest OS
- A grand view of the entire state of guest OS.

How it works




How it works




How it works



How it works



How it works



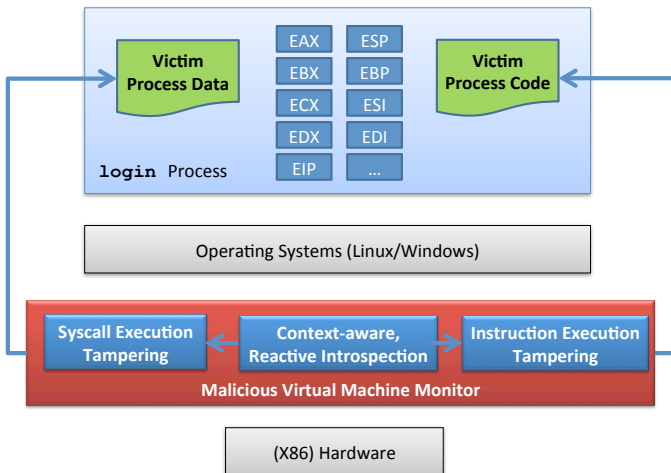
How it works



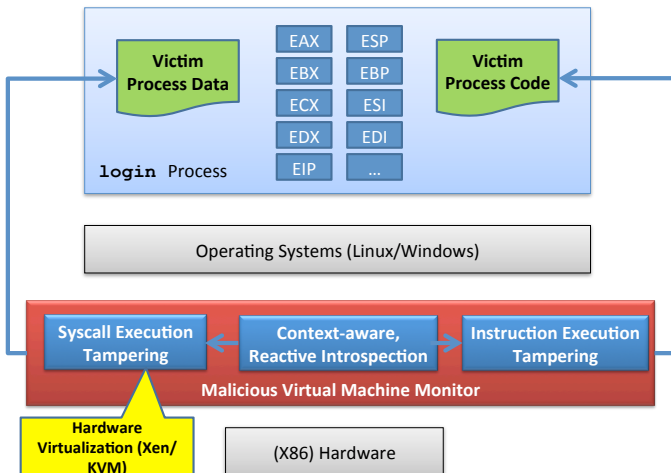
**Malicious Virtual
Machine Monitor**

(X86) Hardware

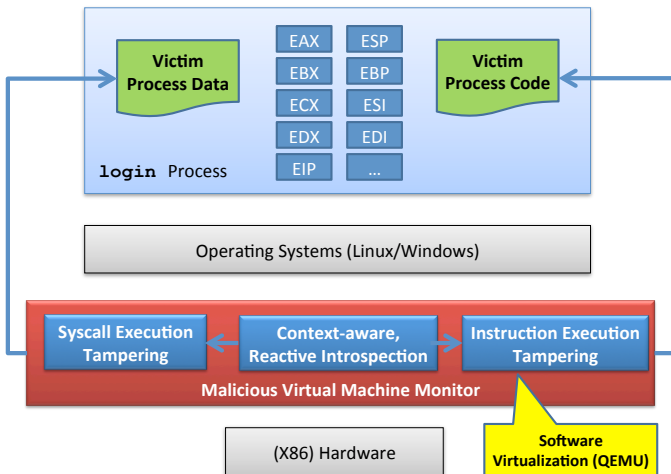
Overview



Using Hardware Virtualization



Using Software Virtualization



Working Example: from instructions perspective

```

if (pw_auth (user_passwd, username, reason, (char *) 0) == 0) {
804a868:  a1 0c 62 05 08          mov     0x805620c,%eax
804a86d:  c7 44 24 0c 00 00 00    movl    $0x0,0xc(%esp)
804a874:  00
804a875:  89 3c 24                mov     %edi,(%esp)
804a878:  89 44 24 08            mov     %eax,0x8(%esp)
804a87c:  a1 48 65 05 08        mov     0x8056548,%eax
804a881:  89 44 24 04            mov     %eax,0x4(%esp)
804a885:  e8 86 87 00 00        call    8053010<pw_auth>
804a88a:  85 c0                  test    %eax,%eax
804a88c:  0f 84 6d fd ff ff      je      804a5ff<main+0x64f>
      goto auth_ok;
}

```

Figure : Binary Code Snippet of the login Program.

Insight-I

Instruction Execution Tampering

- Tampering with Instruction Opcode
 - 804a88c:0f 84 (je) → 0f 85 (jne)
- Tampering with Instruction Operand
 - 804a88a:test %eax,%eax → Tampering w/ eax/EFLAGS
- Tampering with both Opcode and Operand
 - 804a885:call 8053010<pw_auth> → mov \$0,%eax

Working Example: from system call perspective

```

1  execve("/bin/login", ["login"], [/* 16 vars */]) = 0
2  uname({sys="Linux", node="ubuntu", ...}) = 0
...
409 open("/etc/passwd", O_RDONLY) = 4
410 fcntl64(4, F_GETFD) = 0
411 fcntl64(4, F_SETFD, FD_CLOEXEC) = 0
412 _llseek(4, 0, [0], SEEK_CUR) = 0
413 fstat64(4, {st_mode=S_IFREG|0644, st_size=952, ...}) = 0
414 mmap2(NULL, 952, PROT_READ, MAP_SHARED, 4, 0) = 0x4021a000
415 _llseek(4, 952, [952], SEEK_SET) = 0
416 munmap(0x4021a000, 952) = 0
417 close(4) = 0
418 open("/etc/shadow", O_RDONLY) = 4
419 fcntl64(4, F_GETFD) = 0
420 fcntl64(4, F_SETFD, FD_CLOEXEC) = 0
421 _llseek(4, 0, [0], SEEK_CUR) = 0
422 fstat64(4, {st_mode=S_IFREG|0640, st_size=657, ...}) = 0
423 mmap2(NULL, 657, PROT_READ, MAP_SHARED, 4, 0) = 0x4021a000
424 _llseek(4, 657, [657], SEEK_SET) = 0
425 munmap(0x4021a000, 657) = 0
426 close(4) = 0
...

```

Figure : System Call Trace Snippet of the `login` Program.

Insight-II

System Call Execution Tampering

- Tampering with Disk-IO Syscall
 - Replacing `/etc/shadow` file when it loads to the memory. Essentially a man-in-the-middle Attack. We can hijack the file `open` syscall and provide an attacker controlled password file
- Tampering with Memory-Map Syscall
 - Tampering with `mmap2` syscall by replacing the memory contents mapped by this syscall (immediately after it finishes) with the password hash values we control.

Insight-II

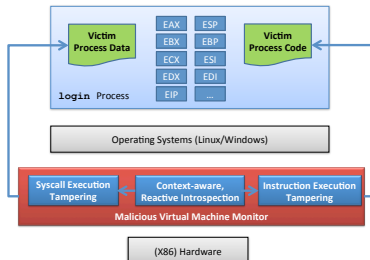
System Call Execution Tampering

- Tampering with Disk-IO Syscall
 - Replacing `/etc/shadow` file when it loads to the memory. Essentially a man-in-the-middle Attack. We can hijack the file `open` syscall and provide an attacker controlled password file
- Tampering with Memory-Map Syscall
 - Tampering with `mmap2` syscall by replacing the memory contents mapped by this syscall (immediately after it finishes) with the password hash values we control.

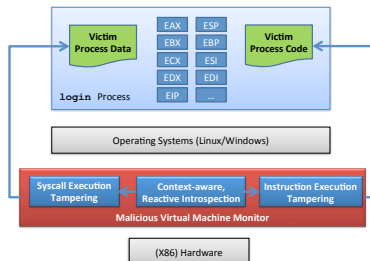
Advantages

- Transparent, can work for many other `login` types of programs
- No binary code reverse engineering

Challenges



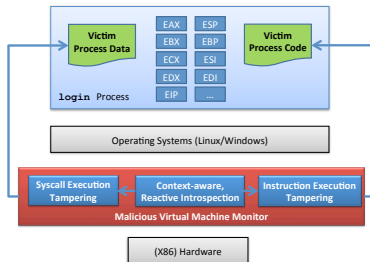
Challenges



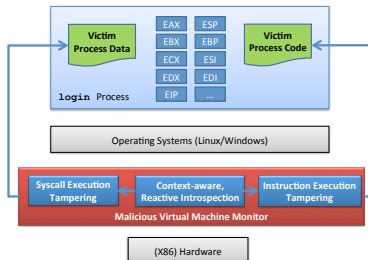
Identifying the “dreaming” context at the VMM layer

- **(C1)** a particular process execution;
- **(C2)** a particular syscall in **C1**;
- **(C3)** a particular instruction in **C1**;
- **(C4)** a particular instruction in **C1** under a particular call stack.

Solutions




Solutions




Context-Aware, reactive introspection

- **Introspection:** a variant of Virtual Machine Introspection [Garfinkel et al, NDSS'03]
- **Reactive:** not a passive, read-only introspection, it is reactive
- **Context-Aware:** context ranges from **C1** to **C4**

Solutions: Designing with Xen/KVM (SYSVMI)




Solutions: Designing with Xen/KVM (SYSVMI)



Execution Context Identification

- **(C1)** – process context: CR3 and code hash of `login`
- **(C2)** – syscall in **C1**:
`sysenter/sysret,int 0x80/iret`

Solutions: Designing with Xen/KVM (SYSVMI)




Execution Context Identification

- **(C1)** – process context: CR3 and code hash of `login`
- **(C2)** – syscall in **C1**:
`sysenter/sysret,int 0x80/iret`


Attack Strategies

- **A1**: Tampering with Instruction Code.
- **A2**: Tampering with Syscall Arguments and Return Values
- **A3**: Tampering with Syscall Produced Data
- **A4**: Using IO Virtualization

Solutions: Designing with QEMU (INSTVMI)




Solutions: Designing with QEMU (INSTVMI)



Execution Context Identification

- **(C3)** – instruction execution: Program Counter (PC)
- **(C4)** – call stack: instrumenting `call/ret`

Solutions: Designing with QEMU (INSTVMI)




Execution Context Identification

- **(C3)** – instruction execution: Program Counter (PC)
- **(C4)** – call stack: instrumenting `call/ret`

Attack Strategies

- **A5:** Tampering with Instruction Code at PC Level
- **A6:** Tampering with Instruction Operand
- **A7:** Tampering with Function Call Arguments and Return Values

Implementation




SYSVMI: Using Xen-4.12

Malicious-VMM w/ Xen-4.12	C1~C2	A1	A2	A3	A4	Total
	1,748	17	10	75	45	1,895

- Implementing **A1** to **A4** with only 1,895 LOC in total (a very low cost for attacker).

Implementation



INSTVMI: Using QEMU-1.01

Malicious-VMM w/ QEMU-1.01	C1 ~ C4	A5	A6	A7	Total
	3,513	35	34	25	3,607

- INSTVMI_a ported the SYSVMI implementation (**C1** and **C2**, and **A1** – **A4**) to a most recent QEMU-1.01
- INSTVMI_b implemented the new attacks unique to the software virtualization (**A5** – **A7**) with fine-grained execution context identification (**C3** and **C4**)

Overall Result

Target	SYSVMI			INSTVMI _a			INSTVMI _b		
	A1	A2,A3	A4	A1	A2,A3	A4	A5	A6	A7
login	✓	✓	✓	✓	✓	✓	✓	✓	✓
sshd	✓	✓	✓	✓	✓	✓	✓	✓	✓
vsftpd	✓	✓	✓	✓	✓	✓	✓	✓	✓
telnetd	✓	✓	✓	✓	✓	✓	✓	✓	✓

Table : Effectiveness of our virtual machine inception attack against the authentication program. Each ✓ symbols denotes a successful way of incepting the victim software.

Performance Overhead




Figure : Macro-benchmark Evaluation of the Performance Overhead of Our VMI

Performance Overhead

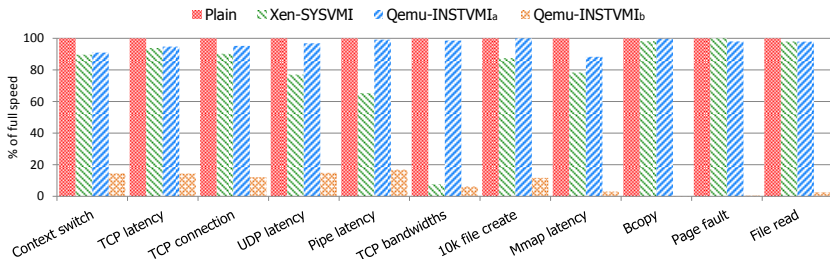
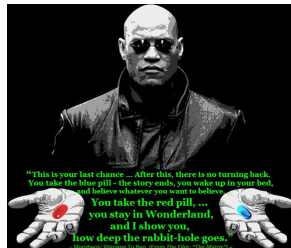


Figure : Micro-benchmark Evaluation of the Performance Overhead of Our VMI

Hardware Virtualization Rootkits

Blue Pill


- The codename for a **rootkit** based on x86 virtualization. [J. Rutkowska, Blackhat'06]
- Trapping a running instance of the OS by starting a **thin** hypervisor
- Vitriol [D. Zov, Blackhat'06].



Key Differences

- Thin vs. Thick Hypervisor
- Our attack explores the relative **freedom** that a malicious virtualization owner has to easily launch stealthy attacks against the virtualized software

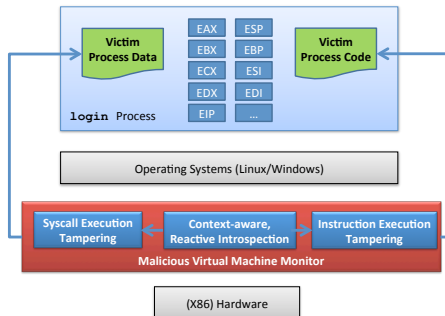
Subvert, SubXen



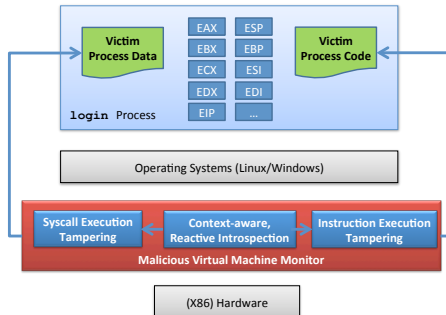
Key Differences

- Subvert aims to infect other's virtualization (to be thin to avoid large footprints)
- Subvert does not bridge the semantic gap, whereas our attack involves a context-aware approach to infer the guest OS semantics.

Summary




Summary



- We design and implement a context-aware, reactive virtual machine to break authentication mechanism.
- Our result indicates that the approach is practical against real-world authentication programs.
- It is useful for both malicious attack and forensics analysis of virtualized systems and software.

Take Away



This attack is so easy

- Small lines of code
- Can be done by anyone, given physical access

Rethink the login design

- login can be executed in a “dreaming” context now
- We have to redesign login to counter against, or at least raise the bar for this attack

Thank you

Questions?

To contact us: {yangchun.fu,zhiqiang.lin,hamlen}@utdallas.edu