

Malware Memory Forensics (MMF) Workshop at ACSAC'14
**Automating Introspection and Forensics
Software Development via Binary Code Reuse**

Zhiqiang Lin

Department of Computer Sciences
The University of Texas at Dallas

December 9th, 2014

Virtualization, Hypervisor, and the Cloud

Windows XP



Linux



Win-7



Virtualization Layer

Hardware Layer

Virtualization, Hypervisor, and the Cloud

Windows XP



Linux



Win-7



..

Virtualization Layer

Hardware Layer

Virtualization (i.e., hypervisor) [Popek and Goldberg, 1974] has pushed our computing paradigm from **multi-tasking** to **multi-OS**.

Virtualization, Hypervisor, and the Cloud

Windows XP



Linux



Win-7



..

Virtualization Layer**Hardware Layer**

Virtualization (i.e., hypervisor) [Popek and Goldberg, 1974] has pushed our computing paradigm from **multi-tasking** to **multi-OS**.

Multiplexing, Isolation, Migration, ...

Virtualization, Hypervisor, and the Cloud

Windows XP



Linux



Win-7



Virtualization Layer

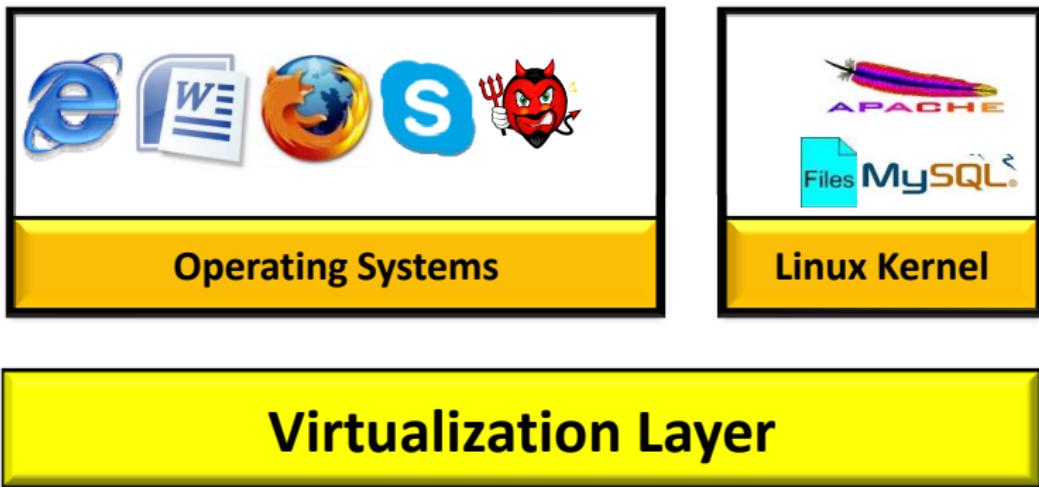
Hardware Layer

Virtualization (i.e., hypervisor) [Popek and Goldberg, 1974] has pushed our computing paradigm from **multi-tasking** to **multi-OS**.

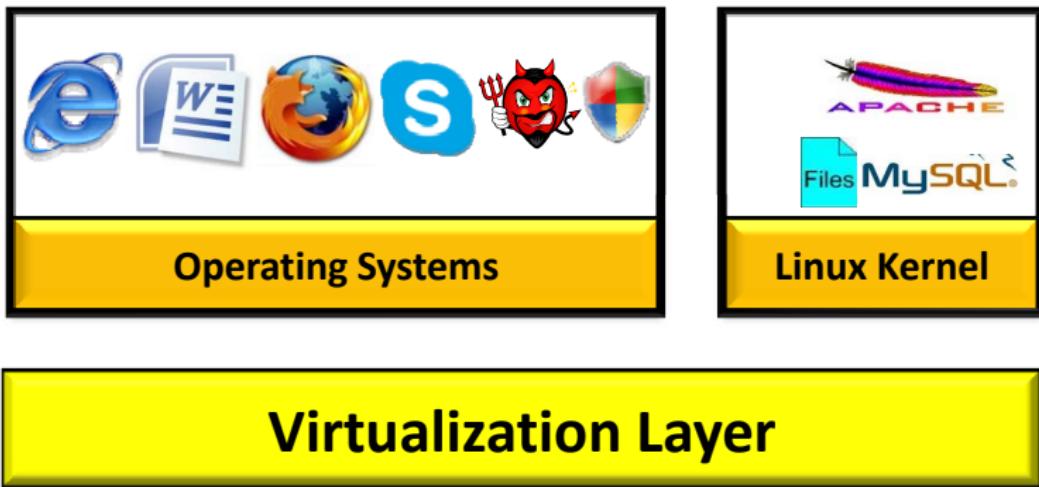
Multiplexing, Isolation, Migration, ...



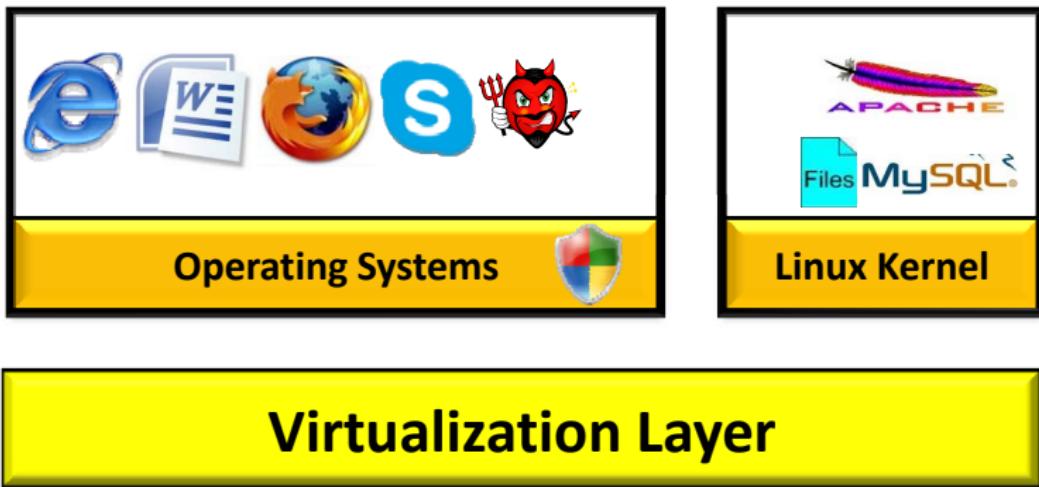
On Anti-Virus Software Evolution



On Anti-Virus Software Evolution



On Anti-Virus Software Evolution



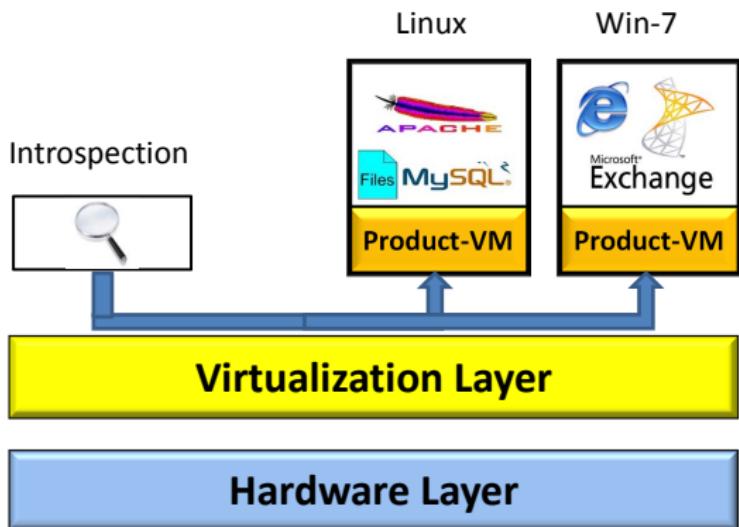
On Anti-Virus Software Evolution



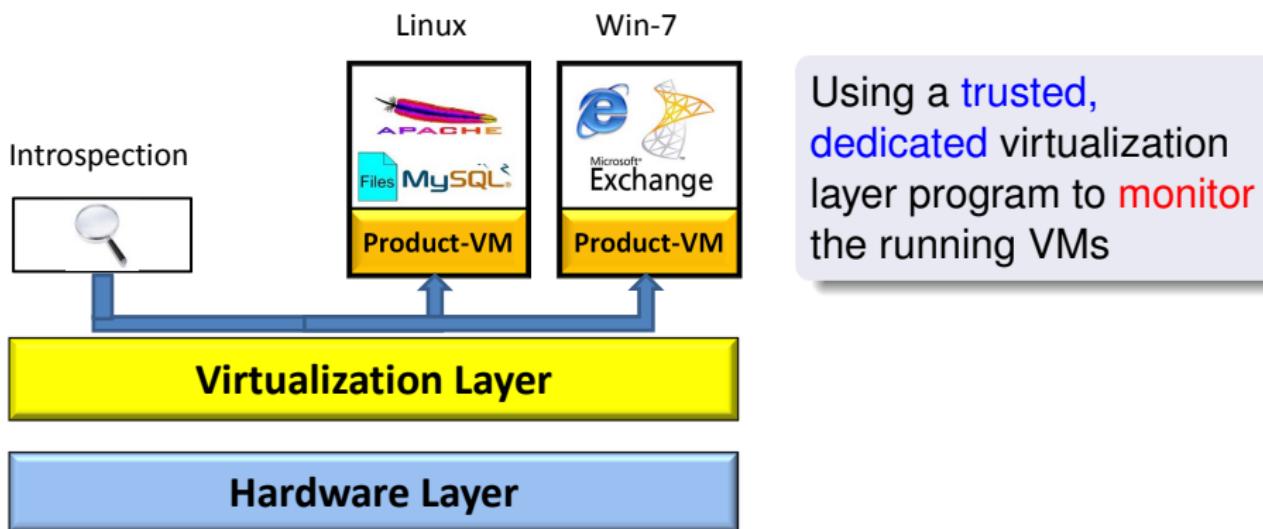
On Anti-Virus Software Evolution



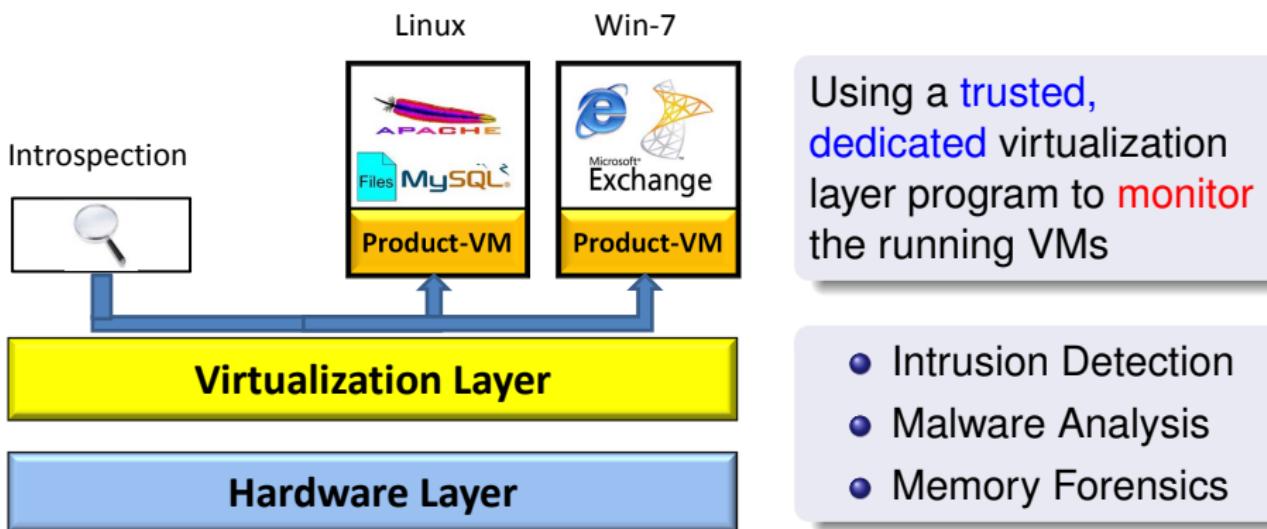
Virtual Machine Introspection (VMI) [Garfinkel et al, NDSS'03]



Virtual Machine Introspection (VMI) [Garfinkel et al, NDSS'03]

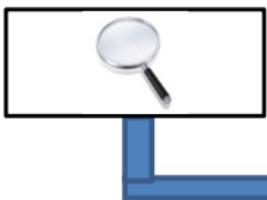


Virtual Machine Introspection (VMI) [Garfinkel et al, NDSS'03]

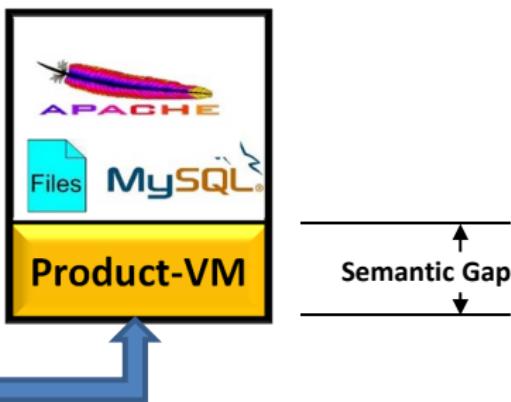


The Semantic Gap in Out-of-VM ([Chen and Noble HotOS'01])

Introspection

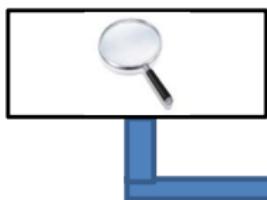


Linux



The Semantic Gap in Out-of-VM ([Chen and Noble HotOS'01])

Introspection



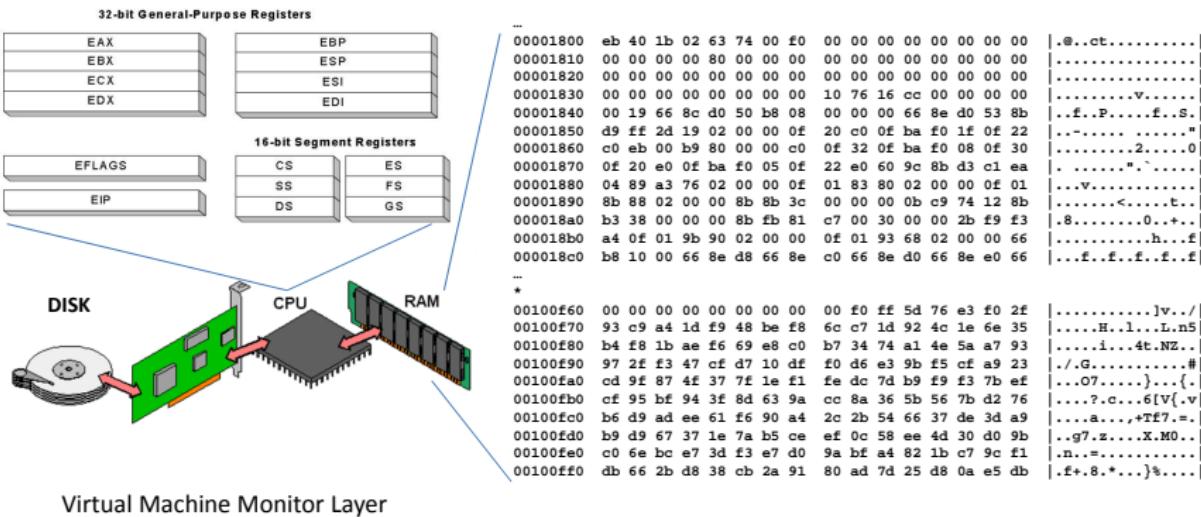
Linux



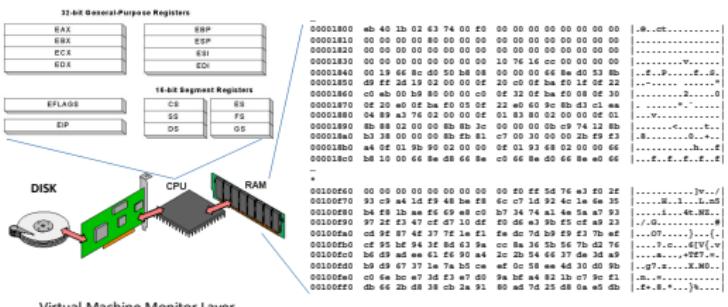
Semantic Gap

- View exposed by Virtual Machine Monitor is at low-level
- There is no abstraction and no APIs
- Need to reconstruct the guest-OS abstraction

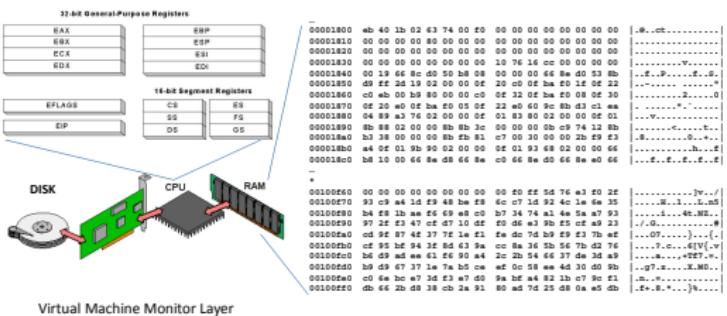
Example: Inspect pids of Guest Memory from VMM



Example: Inspect pids of Guest Memory from VMM



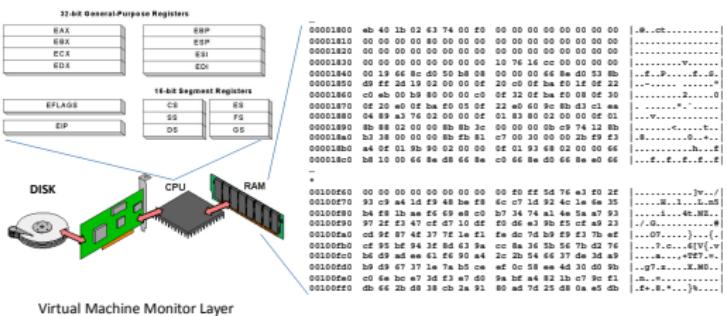
Example: Inspect pids of Guest Memory from VMM



In Kernel 2.6.18

```
struct task_struct {
    ...
    [188] pid_t pid;
    [192] pid_t tgid;
    ...
    [356] uid_t uid;
    [360] uid_t euid;
    [364] uid_t suid;
    [368] uid_t fsuid;
    [372] gid_t gid;
    [376] gid_t egid;
    [380] gid_t sgid;
    [384] gid_t fsgid;
    ...
    [428] char comm[16];
    ...
}
```

Example: Inspect pids of Guest Memory from VMM



- Kernel specific data structure definition
 - Kernel symbols (global variable)
 - Virtual to physical (V2P) translation

In Kernel 2.6.18

```
struct task_struct {
    ...
    [188] pid_t pid;
    [192] pid_t tgid;
    ...
    [356] uid_t uid;
    [360] uid_t euid;
    [364] uid_t suid;
    [368] uid_t fsuid;
    [372] gid_t gid;
    [376] gid_t egid;
    [380] gid_t sgid;
    [384] gid_t fsgid;
    ...
    [428] char comm[16];
}
SIZE: 1408
```

Background
○○○○○

Related Work
●○○○○○○○○

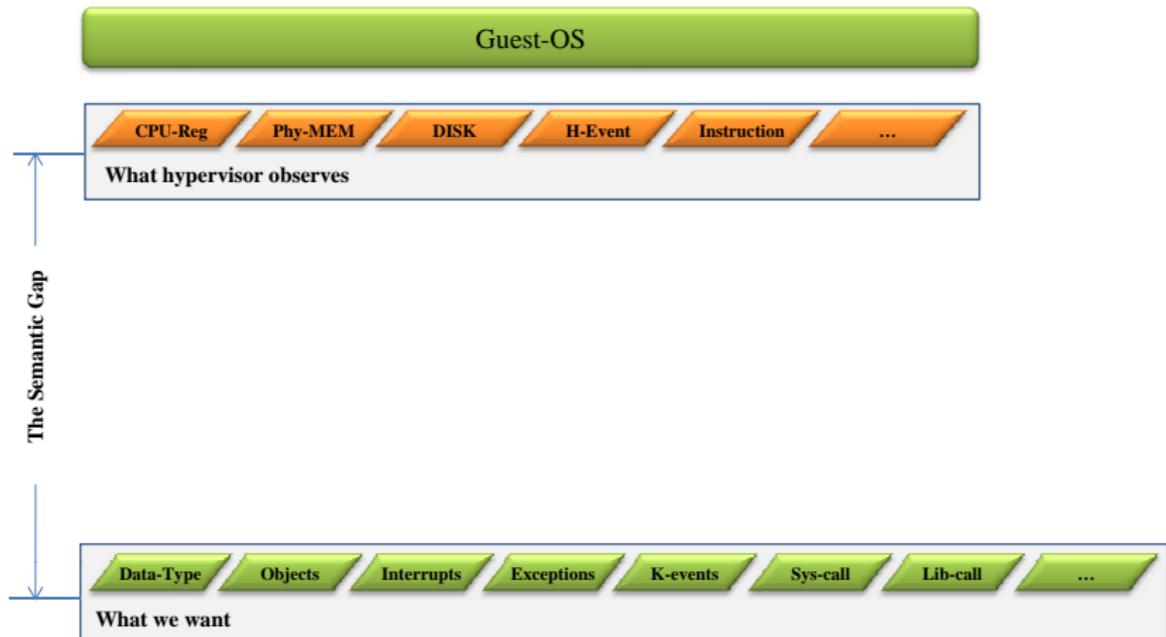
Binary Code Reuse
○○○○○○○

Evaluation
○○○○○○○○○○○○

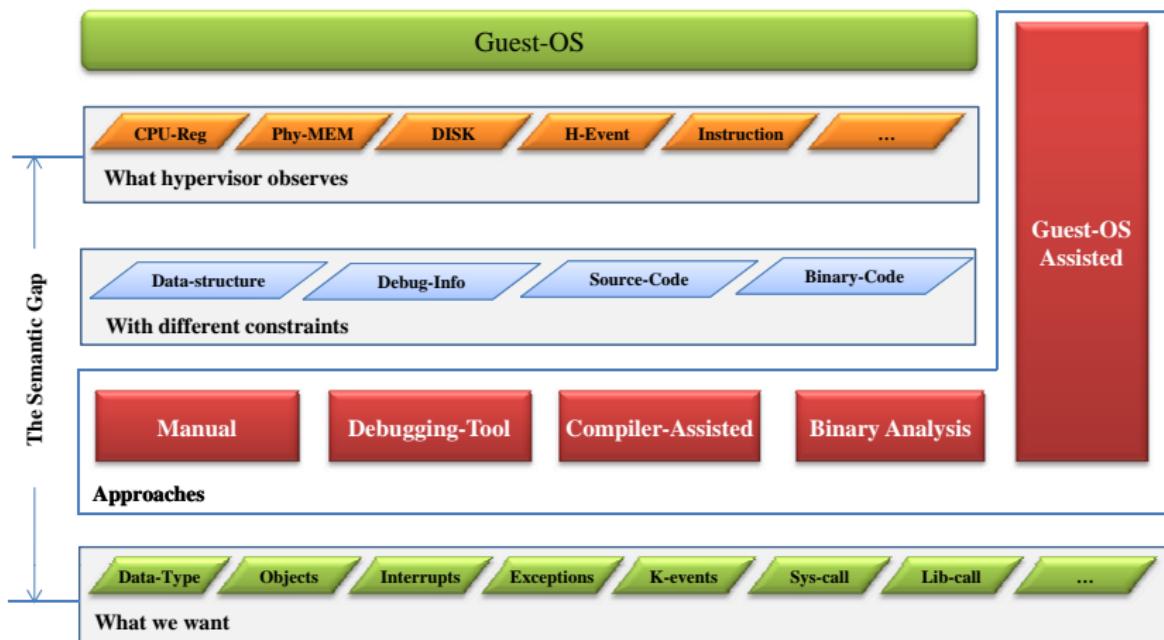
Summary
○○○○○○

How to bridge the semantic gap

How to bridge the semantic gap



How to bridge the semantic gap



State-of-the-art in bridging the Semantic Gap



State-of-the-art in bridging the Semantic Gap



The Semantic Gap
[Chen et al, HotOS'01]

- In HotOS'01, Chen and Noble first raised **the semantic gap problem** in virtualization

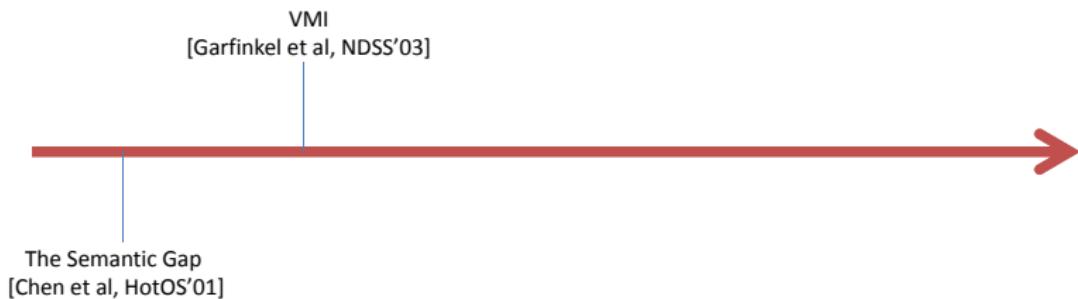
State-of-the-art in bridging the Semantic Gap

The Semantic Gap
[Chen et al, HotOS'01]

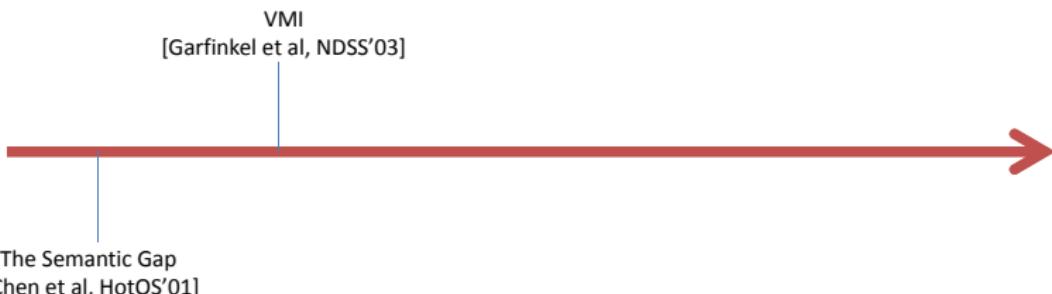
- In HotOS'01, Chen and Noble first raised **the semantic gap problem** in virtualization

“Services in the VM operate **below the abstractions** provided by the guest OS ... This can make it difficult to provide services.”

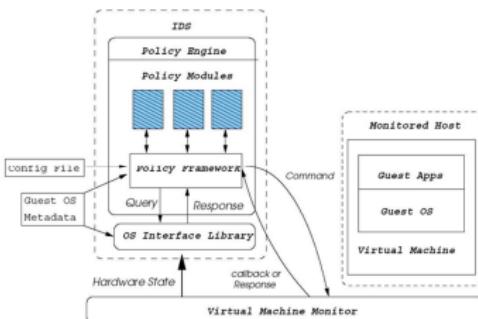
State-of-the-art in bridging the Semantic Gap



State-of-the-art in bridging the Semantic Gap



- In NDSS'03, Garfinkel et al. first proposed VMI, demonstrated for IDS
 - Introspection routine is based on crash utility



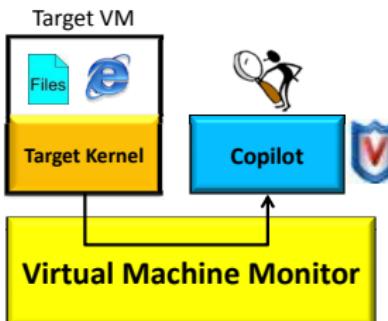
State-of-the-art in bridging the Semantic Gap



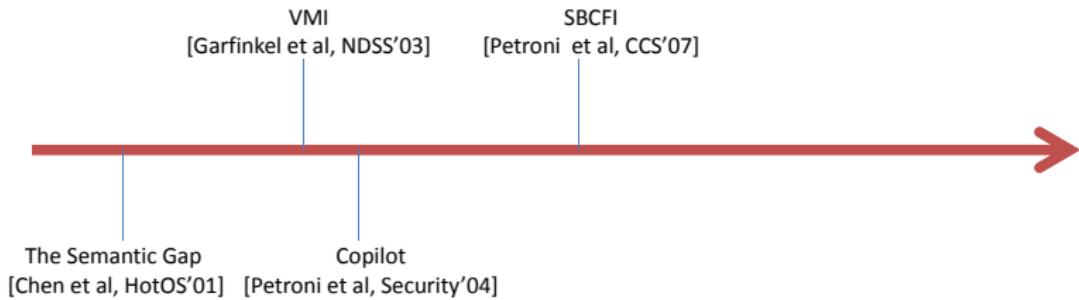
State-of-the-art in bridging the Semantic Gap



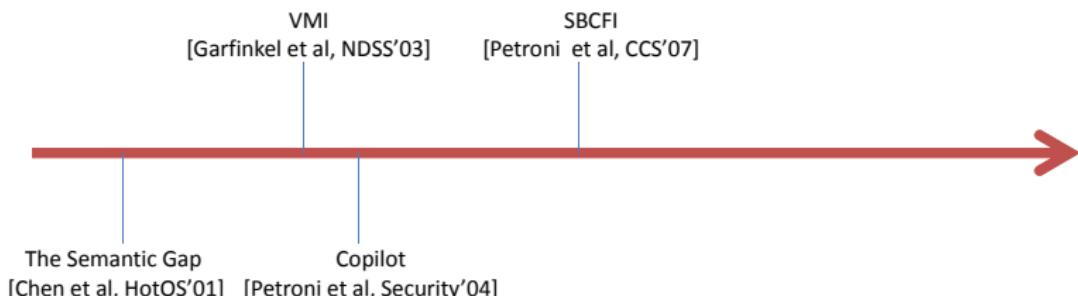
- In USENIX Security'04, Petroni et al. proposed Copilot
 - Introspection routine is based on manually created code



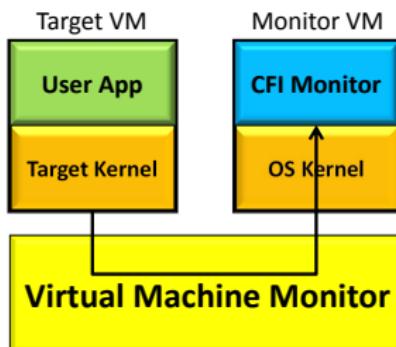
State-of-the-art in bridging the Semantic Gap



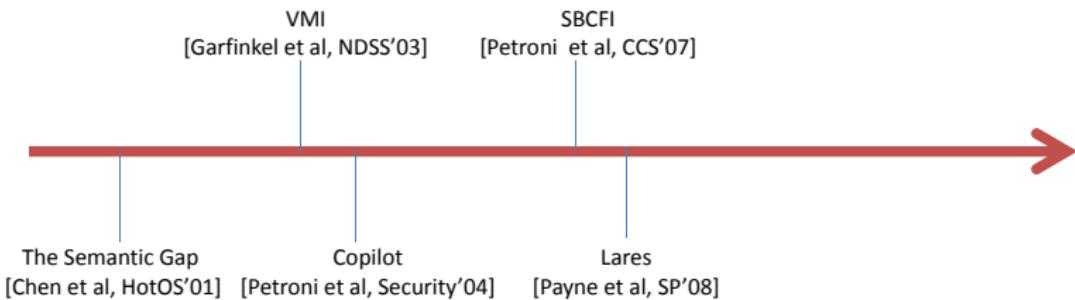
State-of-the-art in bridging the Semantic Gap



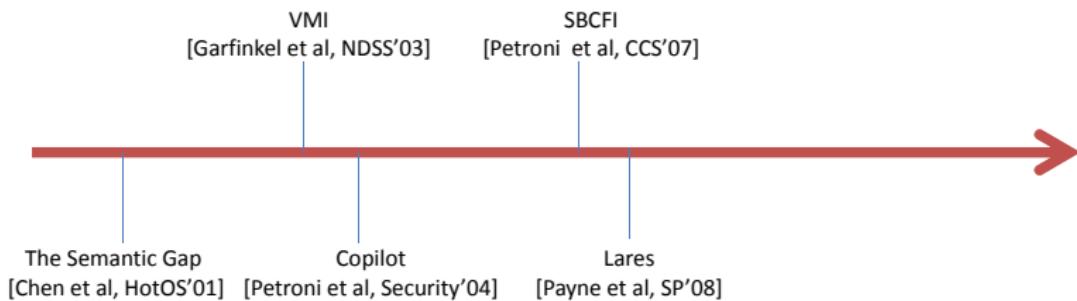
- In CCS'07, Petroni et al. proposed SBCFI
- Introspection routine is based on customized kernel source code



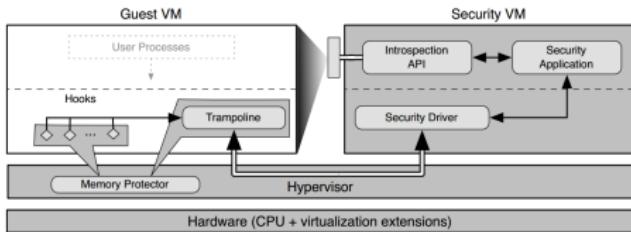
State-of-the-art in bridging the Semantic Gap



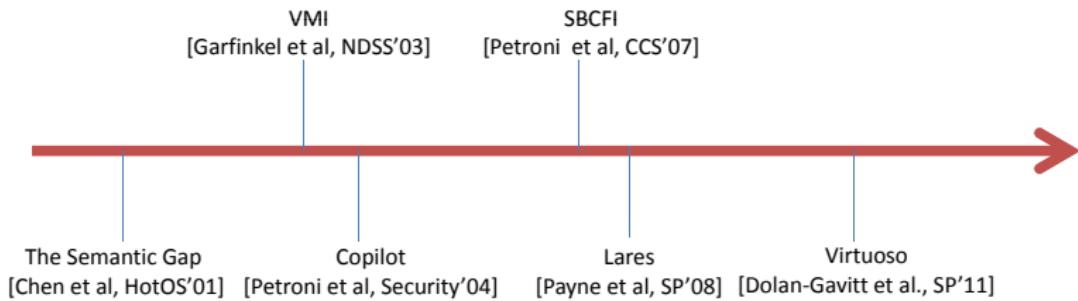
State-of-the-art in bridging the Semantic Gap



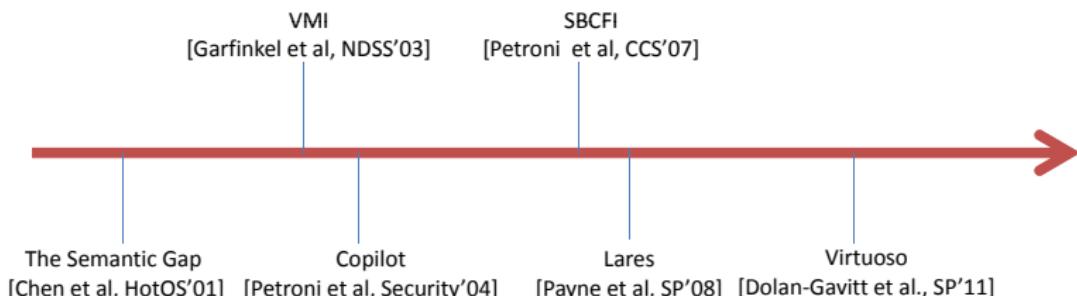
- In SP'08, Payne et al. proposed Lares
- Introspection routine is placed inside the guest OS with special help



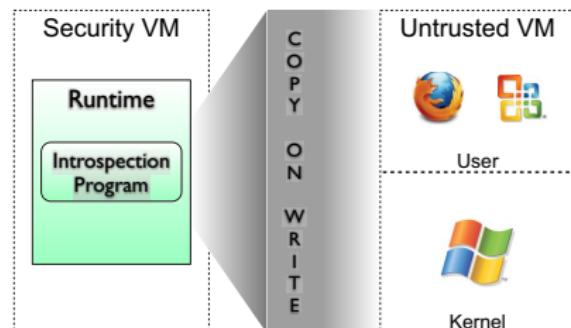
State-of-the-art in bridging the Semantic Gap



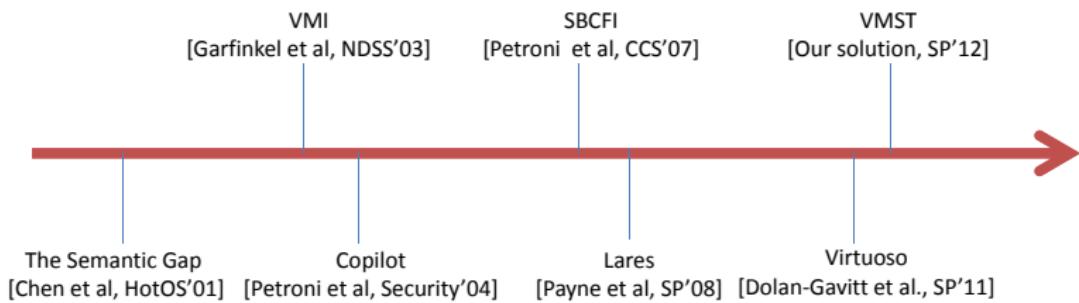
State-of-the-art in bridging the Semantic Gap



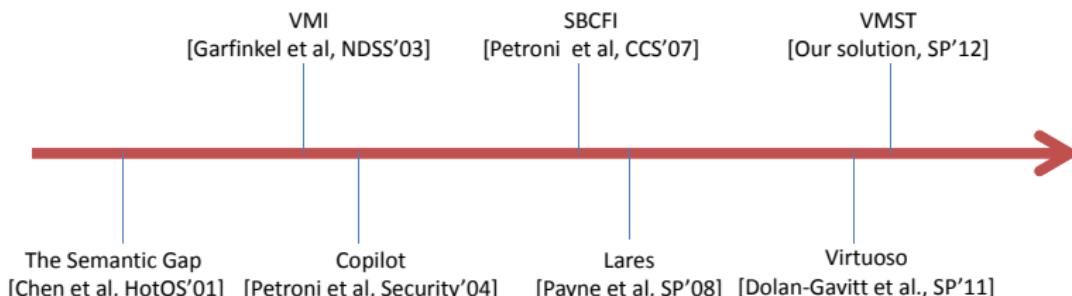
- In SP'11, Dolan-Gavitt et al. proposed Virtuoso
- Introspection routine is based on the trained user level and kernel level code



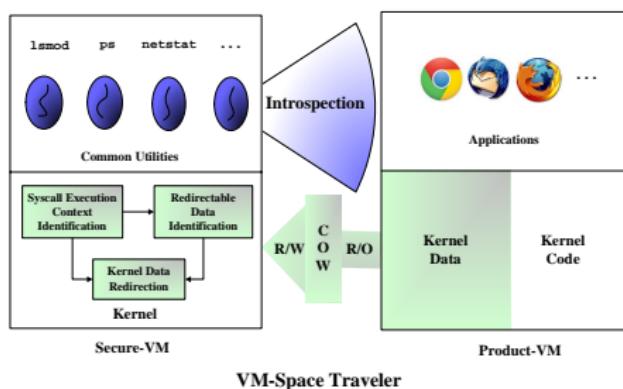
State-of-the-art in bridging the Semantic Gap



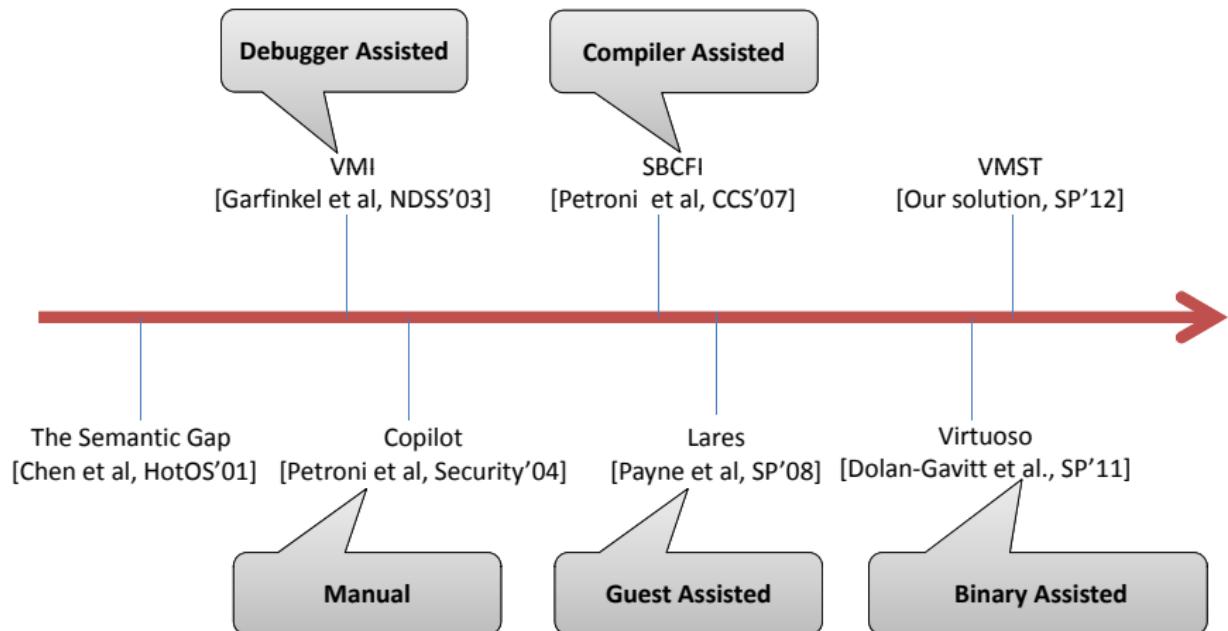
State-of-the-art in bridging the Semantic Gap



- In SP'12, we propose VM space traveling.
- Introspection routine is automatically generated from the **native** user level and kernel level code

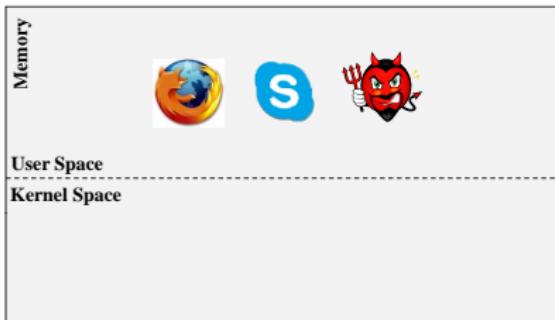


Approach Evolution



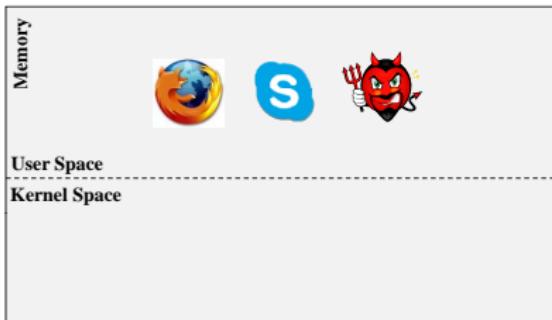
Two Binary Code Reuse Based Approaches

Guest VM (GVM)



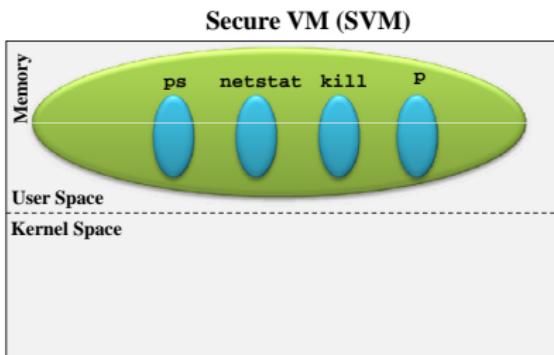
Two Binary Code Reuse Based Approaches

Guest VM (GVM)



Hypervisor

Two Binary Code Reuse Based Approaches



Hypervisor

Using a trusted sibling VM to introspect the running guest VM.

- ① Redirect kernel data [SP'12, VEE'13, NDSS'14]
- ② Redirect system call execution [USENIX ATC'14]

Approach-I: Redirect Kernel Data [SP'12]

In-VM getpid Program

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main()
5 {
6     printf("pid=%d\n",getpid());
7     return 0;
8 }
```

```
1 execve("./getpid",..)          = 0
2 brk(0)                         = 0x83b8000
3 access("/etc/ld.so.nohwcap",..) = -1
...
23 getpid()                      = 13849
...
26 write(1, "pid=13849\n", 10)   = 10
27 exit_group(0)                 = ?
```

Approach-I: Redirect Kernel Data [SP'12]

In-VM getpid Program

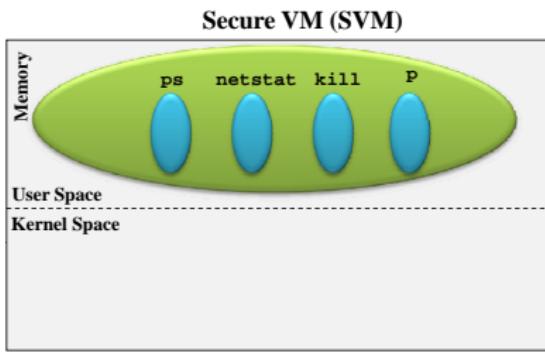
```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main()
5 {
6     printf("pid=%d\n",getpid());
7     return 0;
8 }
```

```
1 execve("./getpid",..)          = 0
2 brk(0)                         = 0x83b8000
3 access("/etc/ld.so.nohwcap",..) = -1
...
23 getpid()                      = 13849
...
26 write(1, "pid=13849\n", 10)   = 10
27 exit_group(0)                 = ?
```

Key Insight

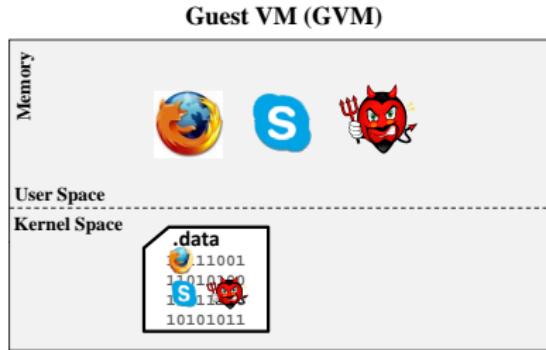
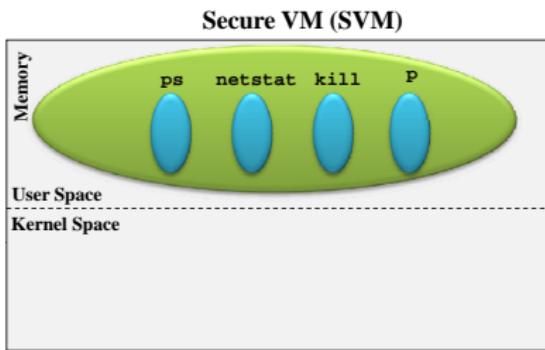
- Reuse the execution context of a native process in SVM.
- When syscall `getpid` executed, redirects the data of interest from the guest VM.

Approach-I: Redirect Kernel Data [SP'12]



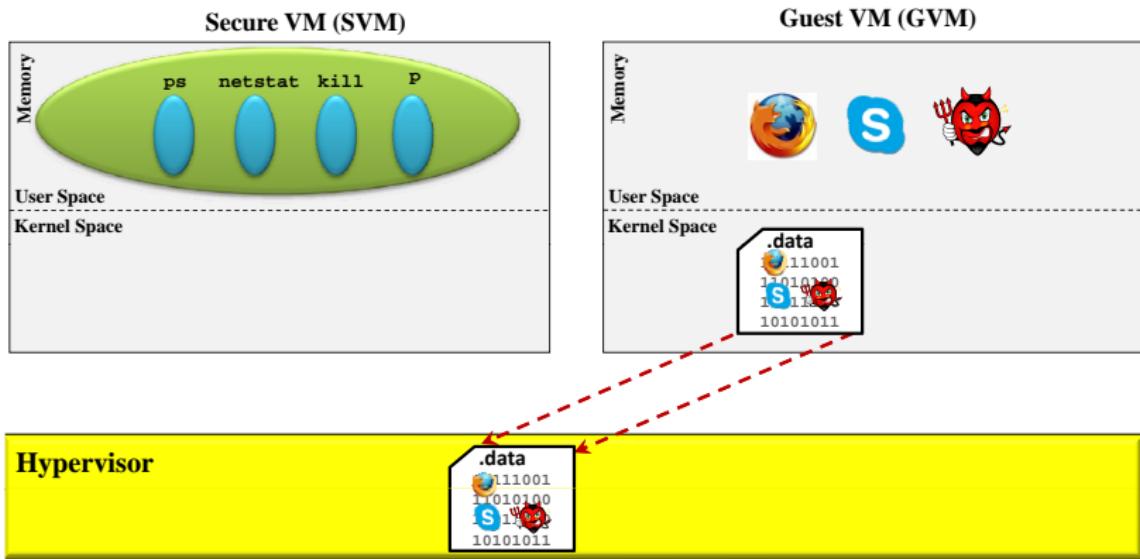
Hypervisor

Approach-I: Redirect Kernel Data [SP'12]

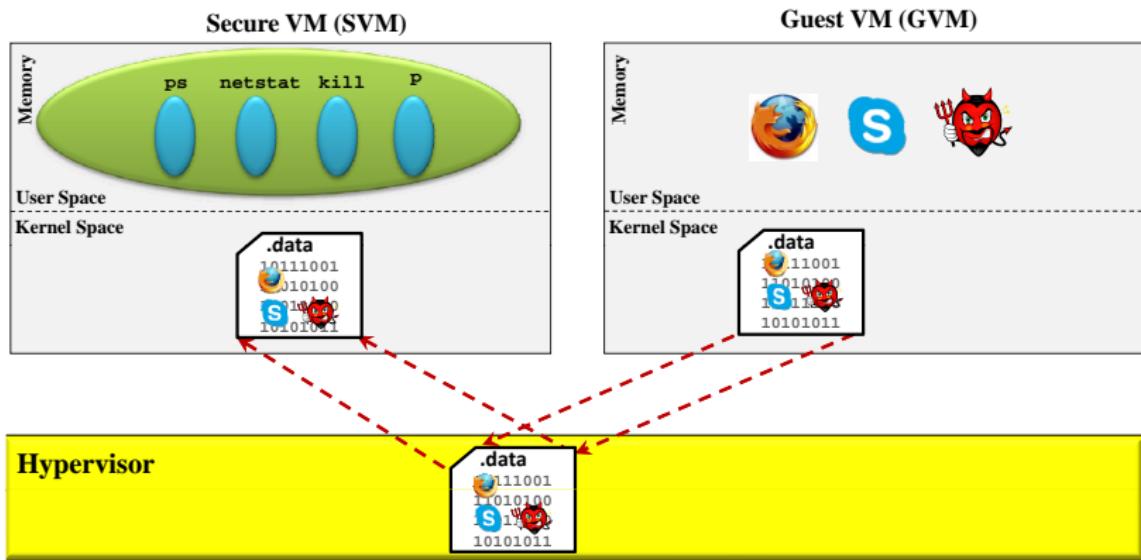


Hypervisor

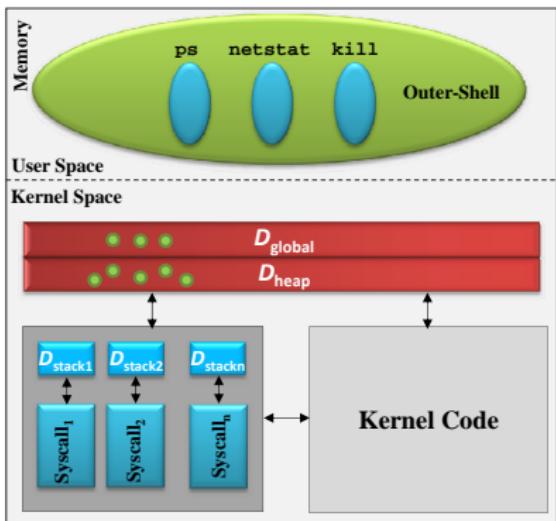
Approach-I: Redirect Kernel Data [SP'12]



Approach-I: Redirect Kernel Data [SP'12]

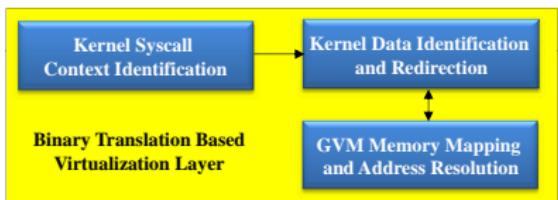
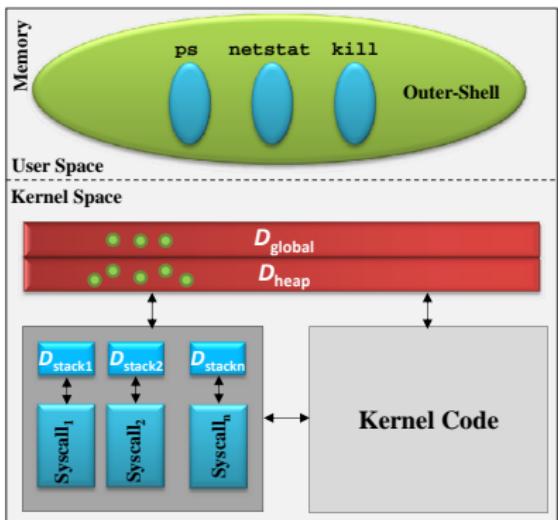


Approach-I: Design & Implementation [SP'12]



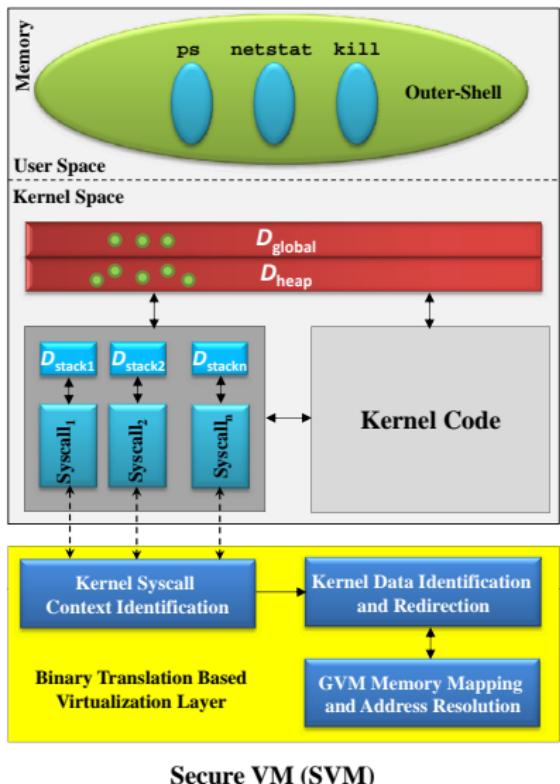
Secure VM (SVM)

Approach-I: Design & Implementation [SP'12]

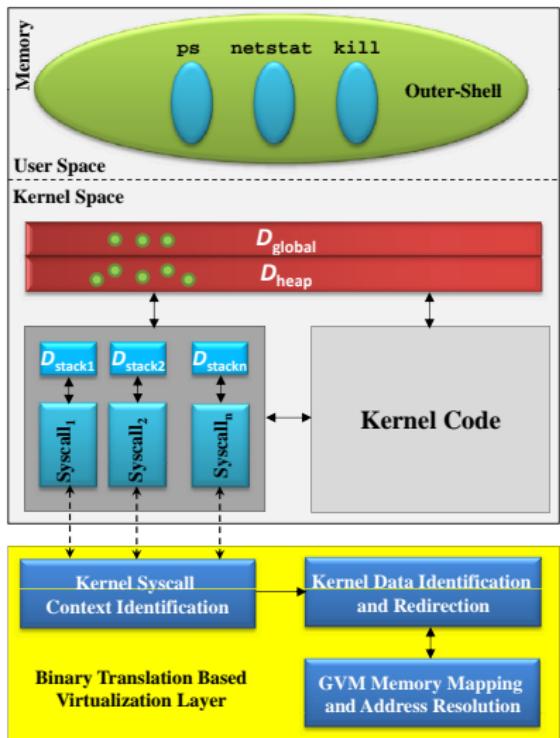


Secure VM (SVM)

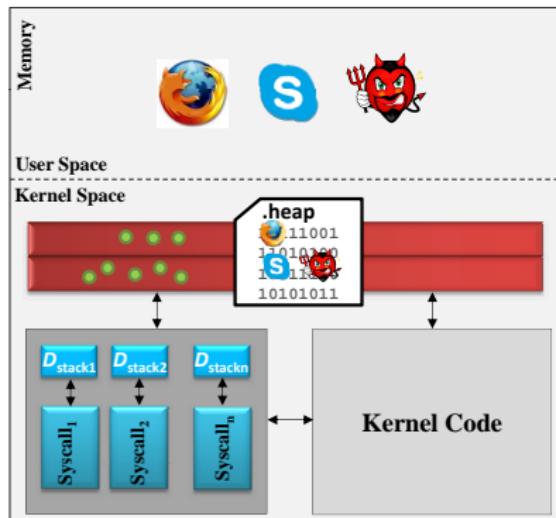
Approach-I: Design & Implementation [SP'12]



Approach-I: Design & Implementation [SP'12]



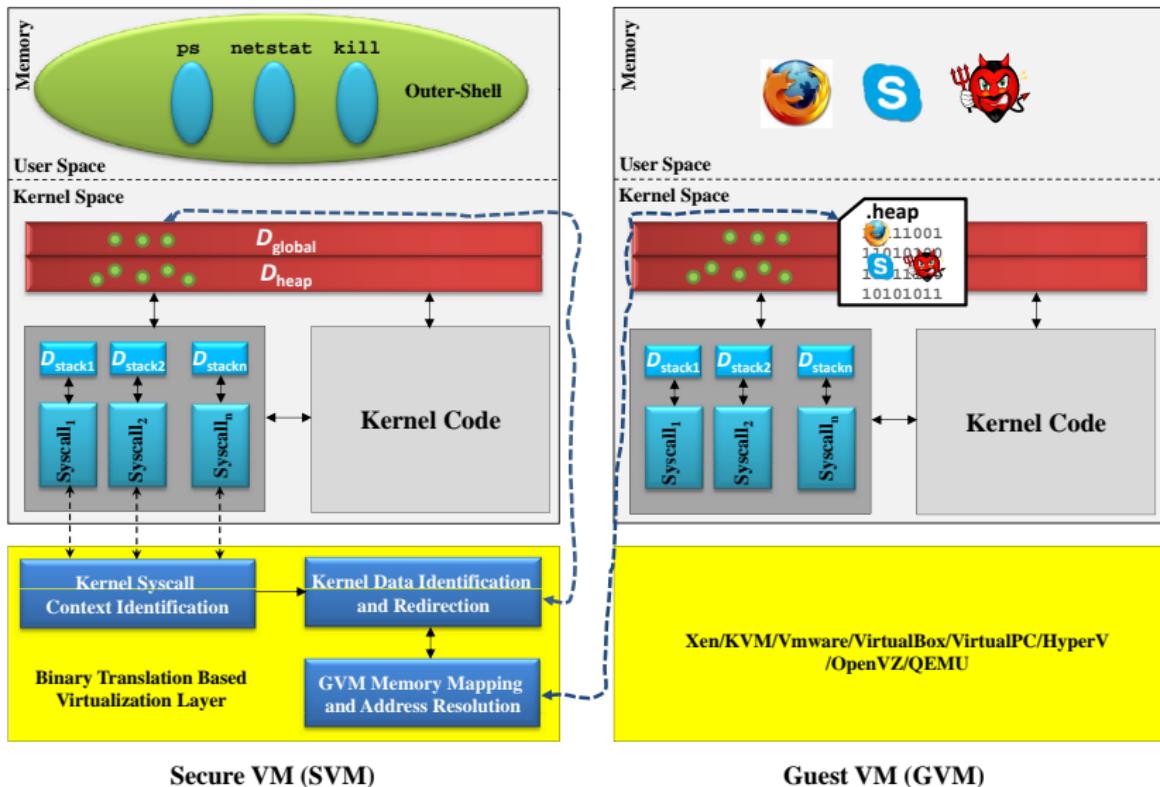
Secure VM (SVM)



Xen/KVM/Vmware/VirtualBox/VirtualPC/HyperV
/OpenVZ/QEMU

Guest VM (GVM)

Approach-I: Design & Implementation [SP'12]



Approach-II: Redirect System call Execution [ATC'14]

In-VM getpid Program

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main()
5 {
6     printf("pid=%d\n",getpid());
7     return 0;
8 }
```

```
1 execve("./getpid",...)          = 0
2 brk(0)                          = 0x83b8000
3 access("/etc/ld.so.nohwcap",..) = -1
...
23 getpid()                      = 13849
...
26 write(1, "pid=13849\n", 10)   = 10
27 exit_group(0)                 = ?
```

Approach-II: Redirect System call Execution [ATC'14]

In-VM getpid Program

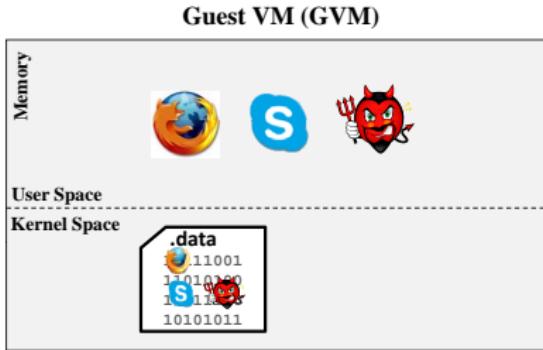
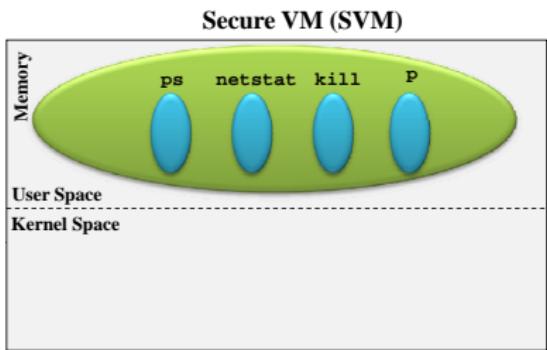
```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main()
5 {
6     printf("pid=%d\n",getpid());
7     return 0;
8 }
```

```
1 execve("./getpid",...)          = 0
2 brk(0)                          = 0x83b8000
3 access("/etc/ld.so.nohwcap",..) = -1
...
23 getpid()                      = 13849
...
26 write(1, "pid=13849\n", 10)   = 10
27 exit_group(0)                 = ?
```

Key Insight

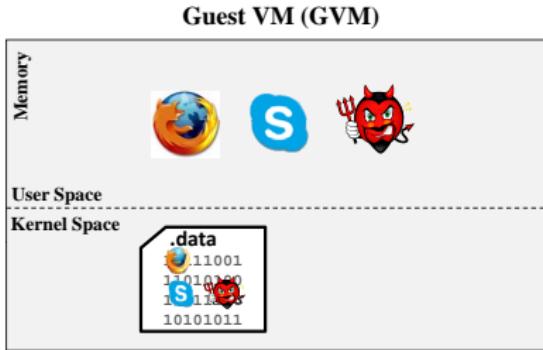
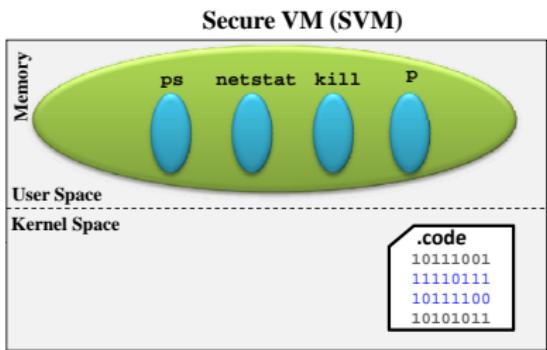
- System call is the only interface to request OS service.
- Pushing the execution of `getpid` system call from SVM to GVM.

Approach-II: Redirect System call Execution [ATC'14]



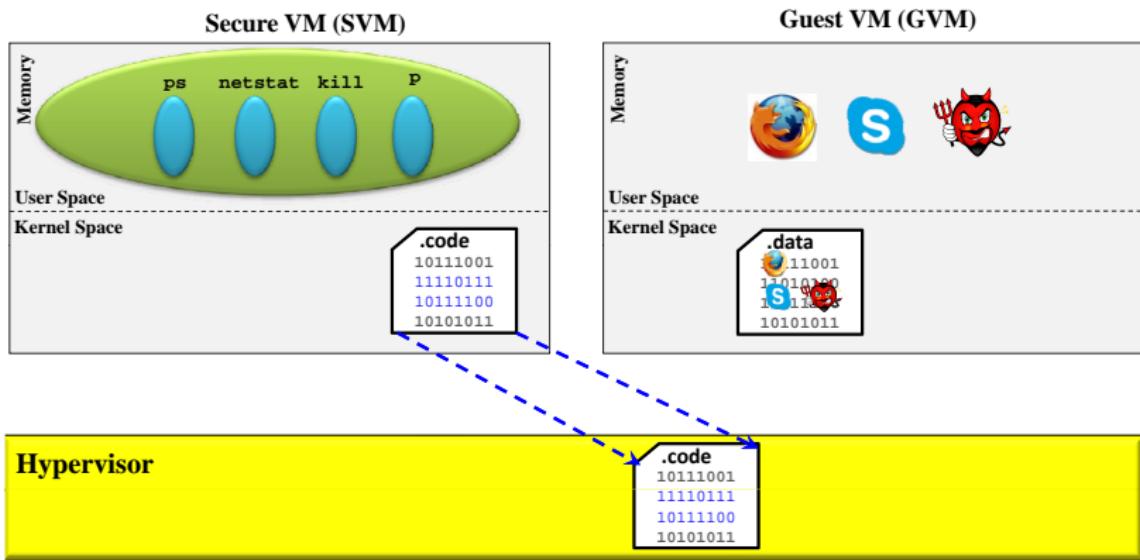
Hypervisor

Approach-II: Redirect System call Execution [ATC'14]

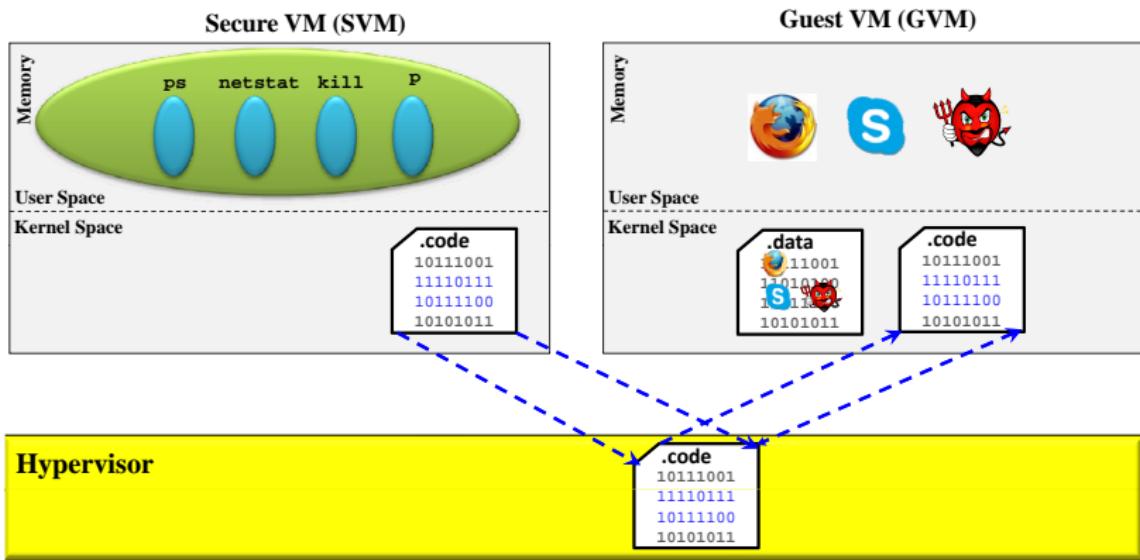


Hypervisor

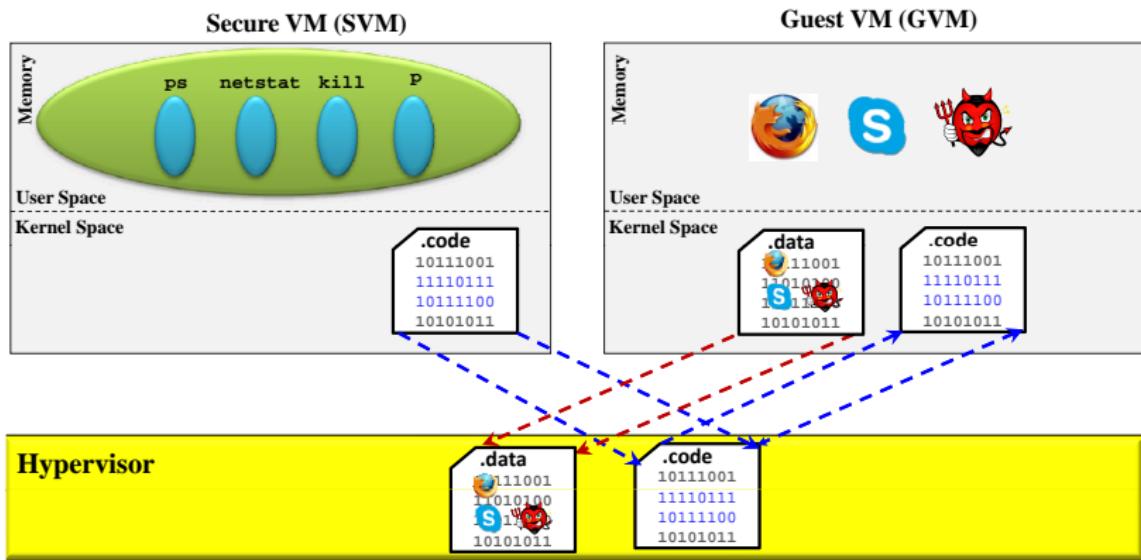
Approach-II: Redirect System call Execution [ATC'14]



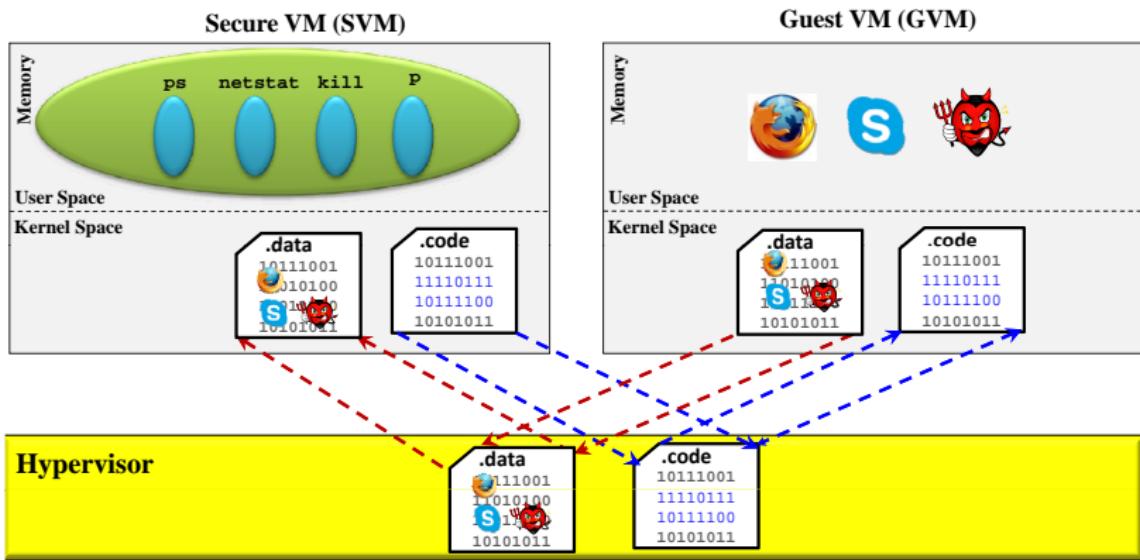
Approach-II: Redirect System call Execution [ATC'14]



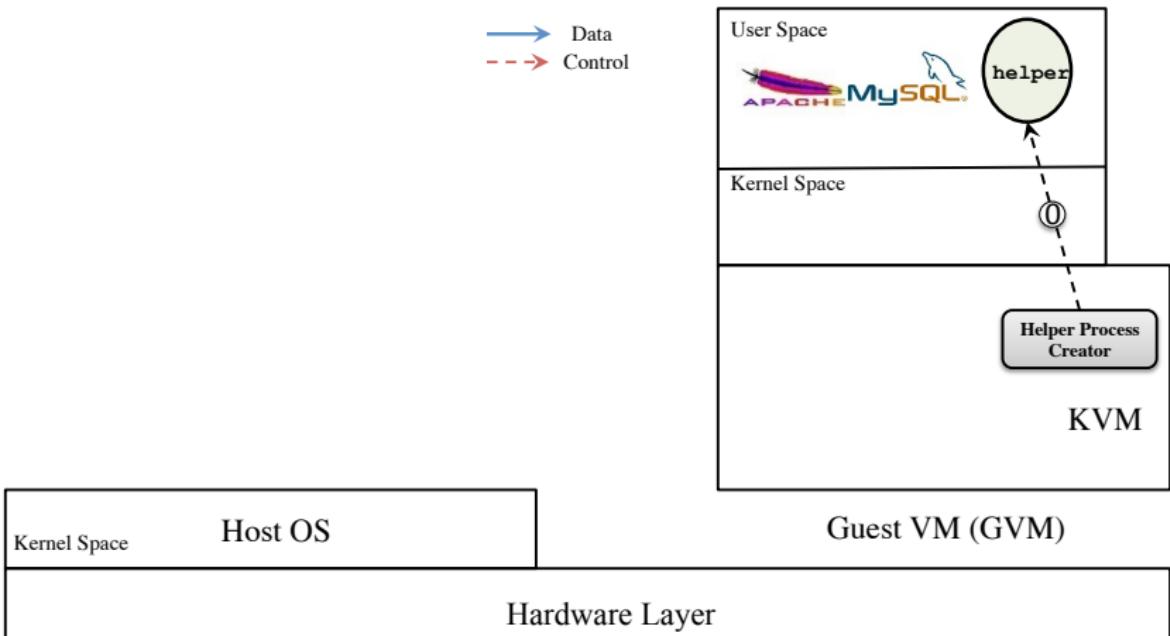
Approach-II: Redirect System call Execution [ATC'14]



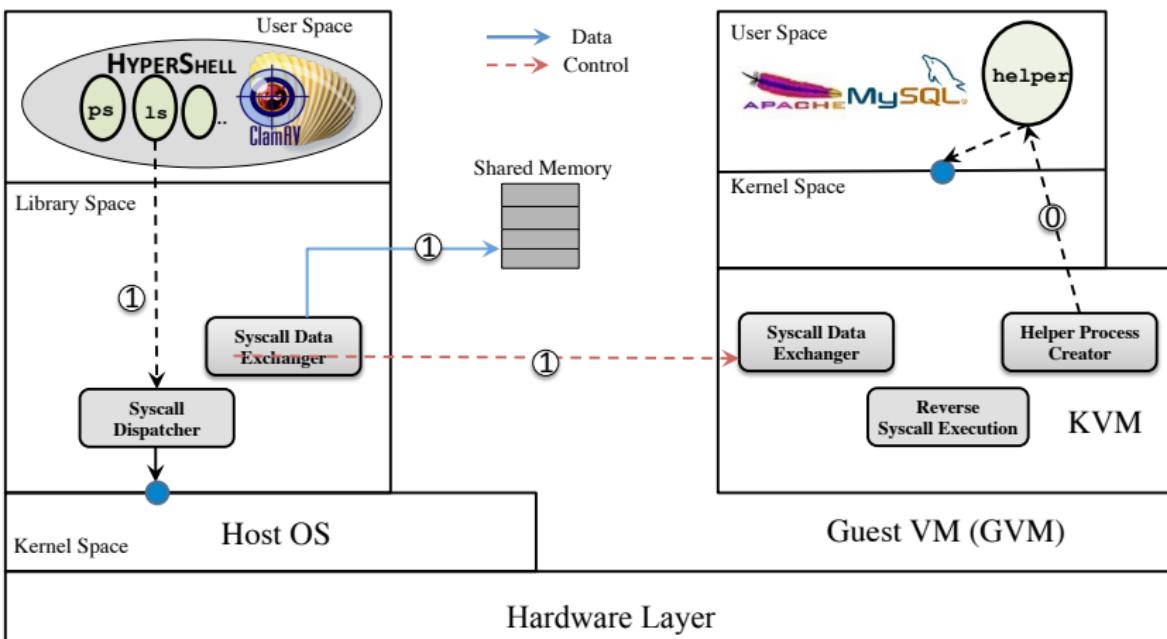
Approach-II: Redirect System call Execution [ATC'14]



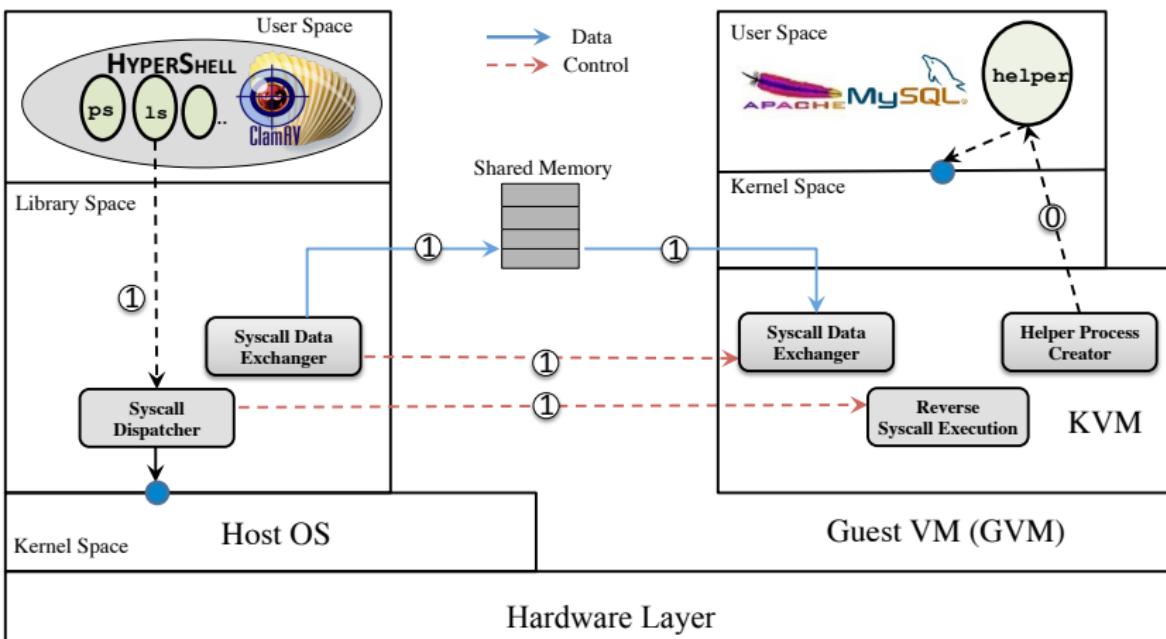
Approach-II: Design & Implementation [ATC'14]



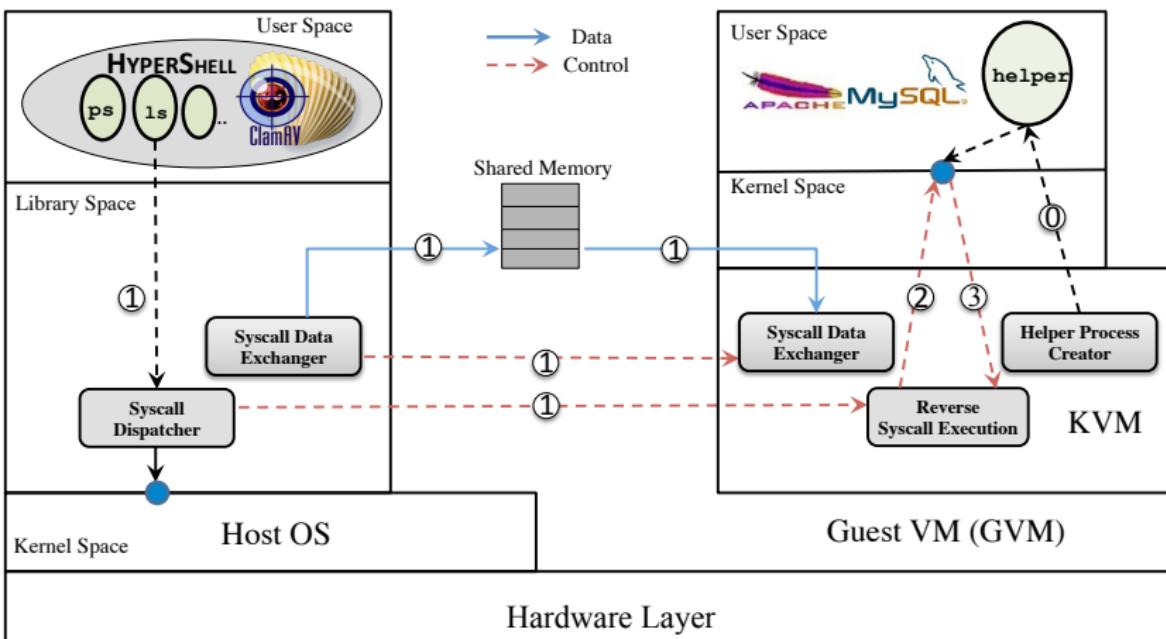
Approach-II: Design & Implementation [ATC'14]



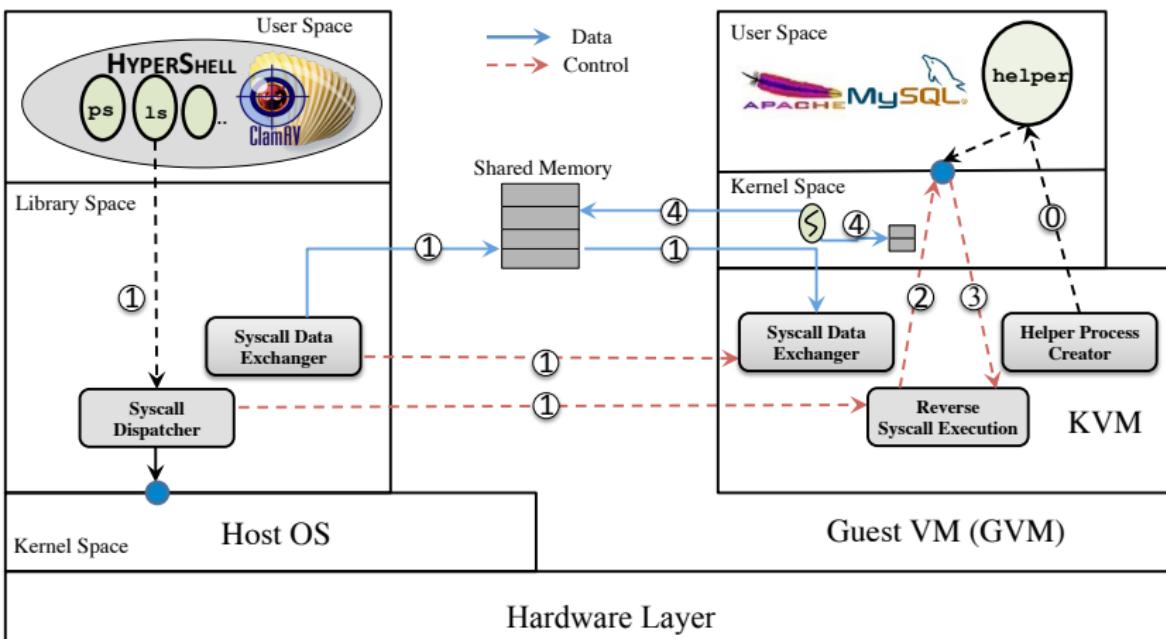
Approach-II: Design & Implementation [ATC'14]



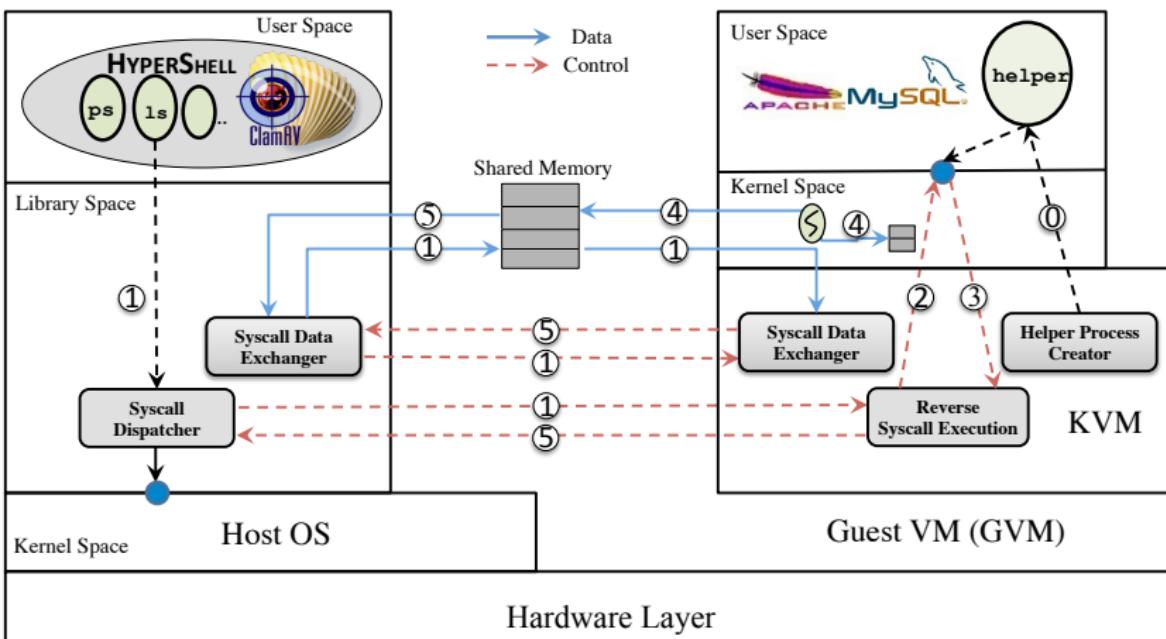
Approach-II: Design & Implementation [ATC'14]



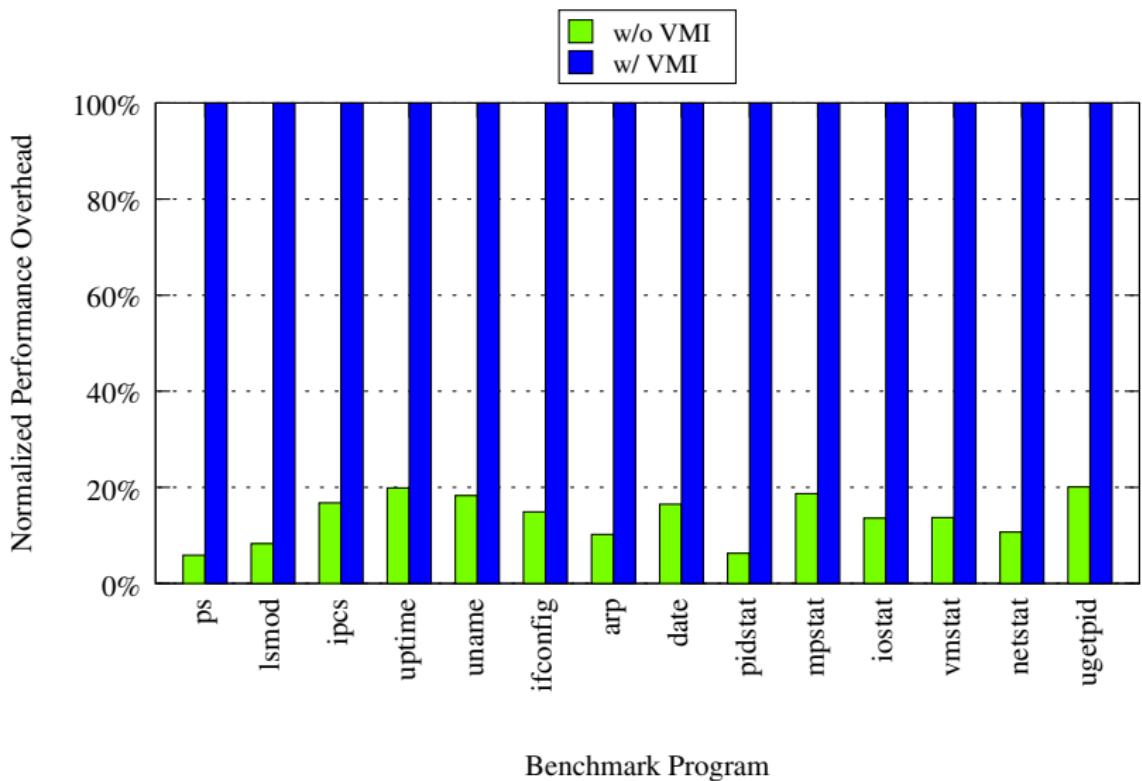
Approach-II: Design & Implementation [ATC'14]



Approach-II: Design & Implementation [ATC'14]



I. Virtual Machine Introspection ([SP'12])

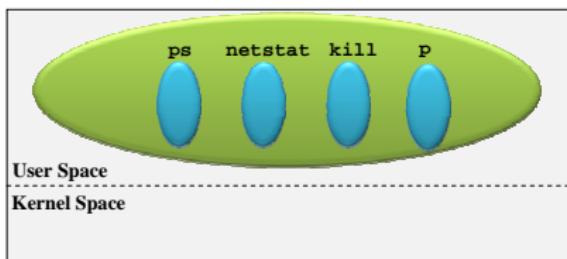


II. Out-of-Box Attack Recovery, Repair

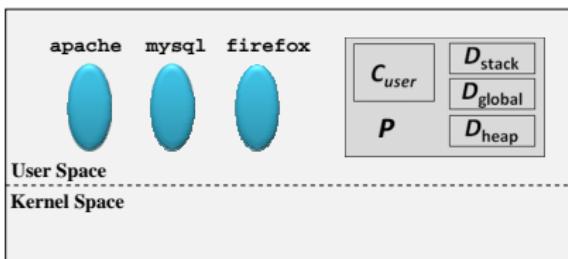
Rootkit	Targeted Function Pointer	Repaired?
adore-2.6	kernel global, heap object	✗
hookswrite	IDT table	✓
int3backdoor	IDT table	✓
kbdv3	syscall table	✓
kbeast-v1	syscall table, tcp4_seq_show	✓
mood-nt-2.3	syscall table	✓
override	syscall table	✓
phalanx-b6	syscall table, tcp4_seq_show	✓
rkit-1.01	syscall table	✓
rial	syscall table	✓
suckit-2	IDT table	✓
synapsys-0.4	syscall table	✓

Table : Rootkit Repairing with An Exterior [VEE'13] Tool.

III. Developing Out-of-VM Programs Natively



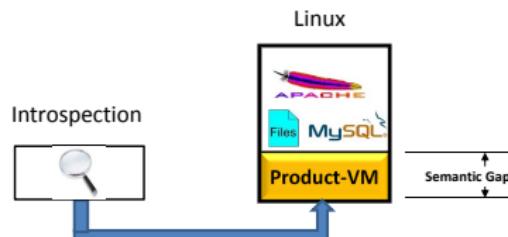
Secure VM (SVM)



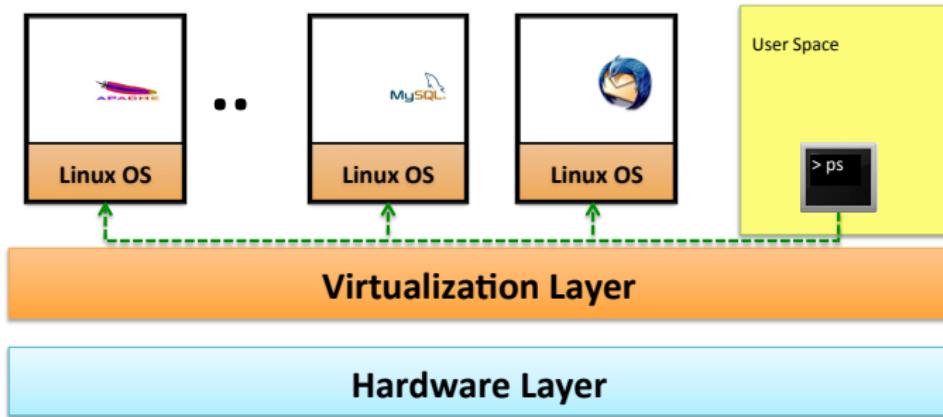
Guest VM (GVM)

In-VM getpid Program

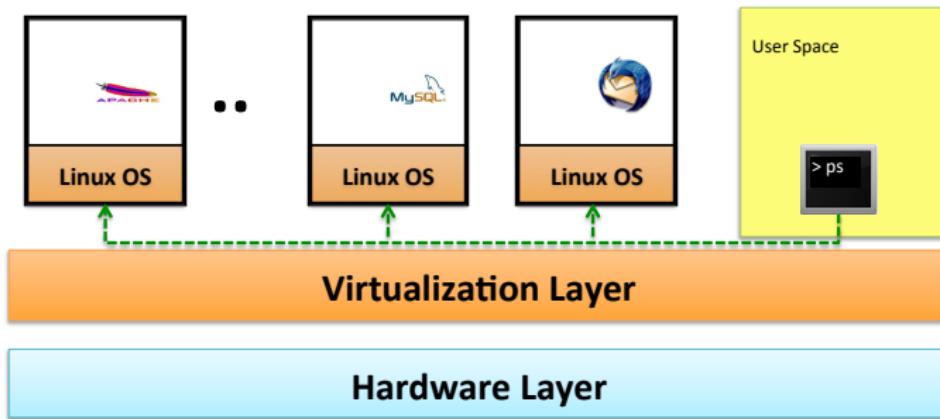
```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main()
5 {
6     printf("pid=%d\n",getpid());
7     return 0;
8 }
```



IV. Writable VMI [VEE'13, ATC'14]



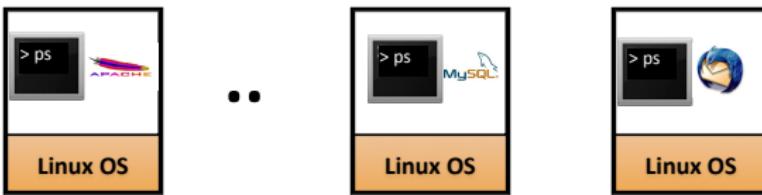
IV. Writable VMI [VEE'13, ATC'14]



Advantages

- Only install the management utilities at hypervisor layer.
- **Automated, uniformed, and centralized management.**

In-VM Management: Existing Approaches



Virtualization Layer

Hardware Layer

In-VM Management: Existing Approaches



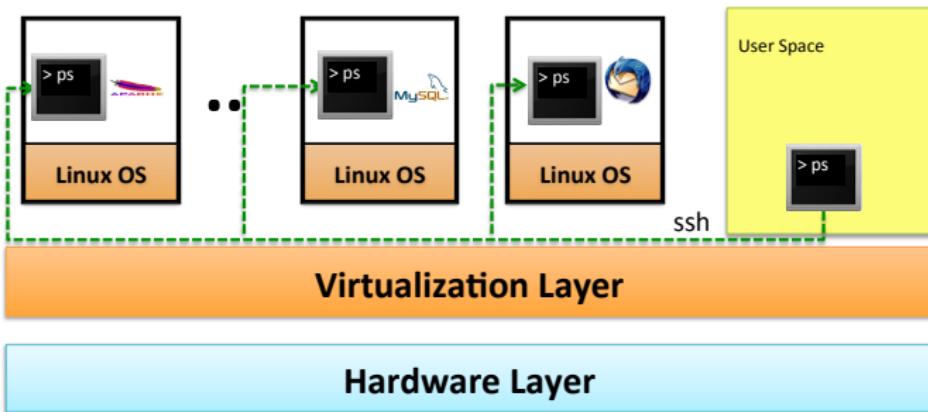
Virtualization Layer

Hardware Layer

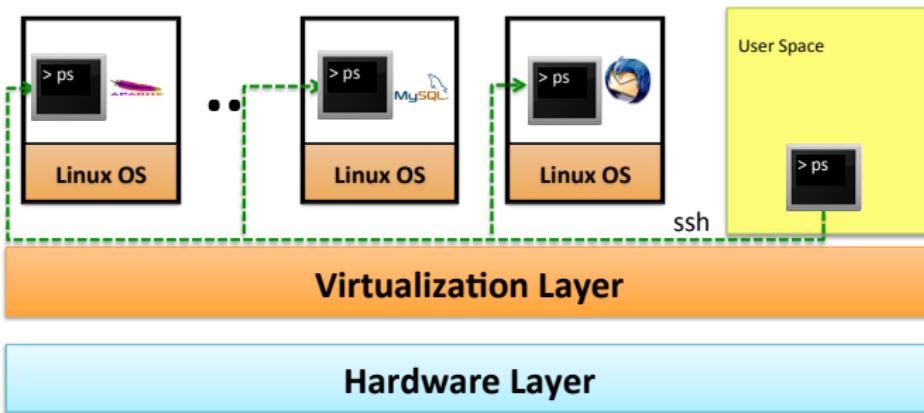
Disadvantages

- Scattered, distributed
- Install, update, and execute in each VM

In-VM Management: Existing Approaches



In-VM Management: Existing Approaches



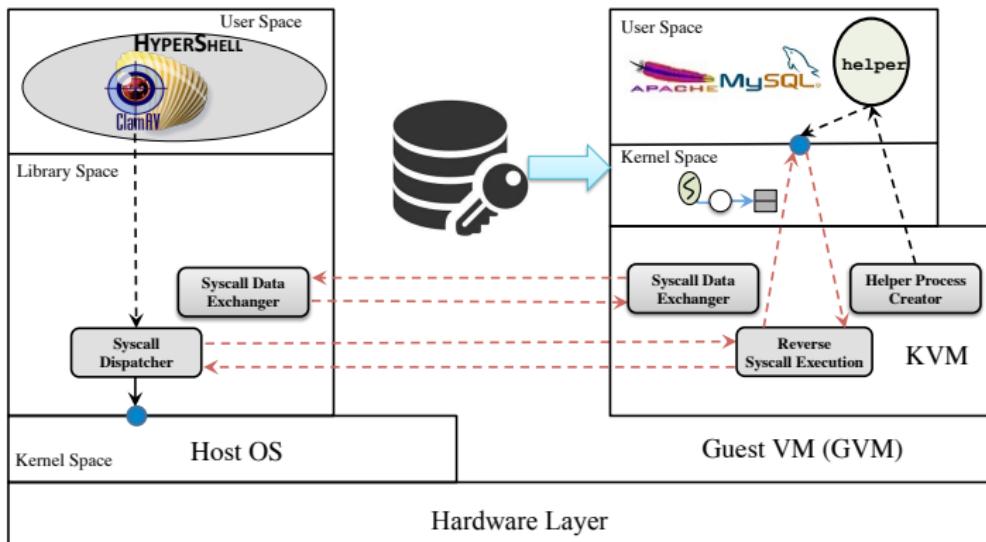
Disadvantages

- Requiring the (admin) login password.
- Requiring install the management utilities in each VM.

Performance Impact: HyperShell [ATC'14]

Process	S	B(ms)	D(ms)	T(X)	date	X	0.11	0.12	1.09	mkdir	✓	0.10	0.19	1.90	
ps	X	1.33	5.42	4.08	w	X	0.95	6.62	6.97	mkfifo	✓	0.10	0.19	1.90	
pidstat	X	1.95	7.56	3.88	hostname	✓	0.04	0.06	1.50	mknod	✓	0.10	0.19	1.90	
nice	✓	0.07	0.11	1.57	groups	✓	0.21	0.62	2.95	mv	✓	0.15	0.31	2.07	
getpid	✓	0.01	0.02	2.00	hostid	✓	0.16	0.56	3.50	rm	✓	0.08	0.15	1.88	
mpstat	X	0.29	0.66	2.28	locale	✓	0.09	0.17	1.89	od	✓	0.12	0.35	2.92	
pstree	X	0.69	6.03	8.74	getconf	✓	0.09	0.34	3.78	cat	✓	0.07	0.18	2.57	
chrt	✓	0.11	0.16	1.45	System Utils					link	✓	0.07	0.13	1.86	
renice	✓	0.11	0.18	1.64	uptime	X	0.07	0.47	6.71	comm	✓	0.08	0.22	2.75	
top	X	504.92	510.85	1.01	sysctl	✓	8.5	42.72	5.03	shred	X	0.72	0.92	1.28	
nproc	✓	0.07	0.26	3.71	arch	✓	0.07	0.11	1.57	truncate	✓	0.07	0.26	3.71	
sleep	✓	1.27	1.28	1.01	dmesg	✓	0.38	0.51	1.34	head	✓	0.07	0.15	2.14	
pgrep	✓	0.89	4.72	5.30	lscpu	✓	0.26	1.21	4.65	vdir	✓	0.63	3.95	6.27	
pkill	✓	0.87	4.33	4.98	mcookie	X	0.29	0.49	1.69	nl	✓	0.08	0.17	2.13	
snice	✓	0.17	0.65	3.82	Disk/Devices					tail	✓	0.08	0.20	2.50	
echo	✓	0.07	0.09	1.29	bblkid	✓	0.14	0.61	4.36	namei	✓	0.07	0.13	1.86	
pwdx	✓	0.05	0.07	1.40	badblocks	✓	0.35	0.44	1.26	whereis	✓	2.05	4.86	2.37	
pmap	✓	0.16	0.36	2.25	lspci	✓	31.40	36.52	1.16	stat	✓	0.27	0.78	2.89	
kill	✓	0.01	0.04	4.00	iostat	✓	0.45	1.04	2.31	readlink	✓	0.07	0.12	1.71	
killall	✓	0.62	3.03	4.89	du	✓	0.11	0.53	4.82	unlink	✓	0.07	0.13	1.86	
Memory	S	B(ms)	D(ms)	T(X)	df	✓	0.16	0.35	2.19	cut	✓	0.08	0.17	2.13	
	free	X	0.04	0.08	2.00	Filesystem					dir	✓	0.07	0.20	2.86
	vmstat	X	0.19	0.33	1.74	sync	✓	8.07	6.53	0.81	mktemp	✓	0.09	0.18	2.00
slabtop	X	0.22	0.36	1.64	getcap	✓	0.04	0.08	2.00	rmdir	✓	0.07	0.13	1.86	
Modules	S	B(ms)	D(ms)	T(X)	lsof	✓	3.31	6.12	1.85	ptx	✓	0.12	0.45	3.75	
	rmmod	✓	0.51	3.14	6.16	pwd	✓	0.07	0.11	1.57	chcon	✓	0.06	0.12	2.00
	modinfo	✓	0.48	1.54	3.21	Files					Network	S	B(ms)	D(ms)	T(X)
lsmod	✓	0.10	0.17	1.70	chggrp	✓	0.19	0.47	2.47	ifconfig	X	0.32	1.15	3.59	
Environment	S	B(ms)	D(ms)	T(X)	chmod	✓	0.07	0.14	2.00	ip	✓	0.10	0.20	2.00	
	who	✓	0.14	0.72	chown	✓	0.19	0.47	2.47	route	✓	138.65	150.32	1.08	
	env	✓	0.07	0.11	cp	✓	0.11	0.27	2.45	ipmaddr	✓	0.13	0.34	2.62	
	printenv	✓	0.07	0.1	uniq	✓	0.09	0.35	3.89	iptunnel	✓	0.09	0.29	3.22	
	whoami	✓	0.19	0.45	file	✓	0.87	1.72	1.98	nameif	✓	0.10	0.21	2.10	
	stty	✓	0.11	0.46	find	✓	0.20	0.58	2.90	netstat	X	0.25	0.37	1.48	
	users	✓	0.09	0.53	grep	✓	0.35	2.14	6.11	arp	✓	0.14	0.24	1.71	
	uname	✓	0.09	0.11	ln	✓	0.08	0.14	1.75	ping	X	15.02	18.2	1.21	
	id	✓	0.26	0.85	ls	✓	0.14	0.27	1.93	Avg.	-	7.27	8.45	2.73	

Disk Introspection: FDE disk virus scanning [ATC'14]



Disk Introspection: FDE disk virus scanning [ATC'14]



- 1. Encrypted by dm-crypt
- 2. 101,415 files
- 3. 1336.09 megabytes in size

Disk Introspection: FDE disk virus scanning [ATC'14]



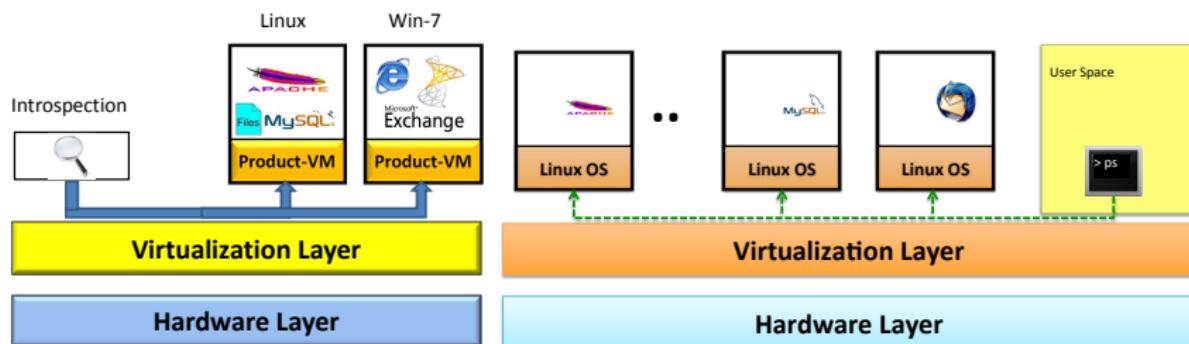
- 1. Encrypted by dm-crypt
- 2. 101,415 files
- 3. 1336.09 megabytes in size

Clamav successfully detect two viruses!!

Comparison with the most related work

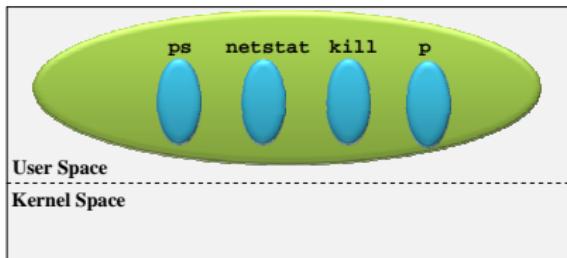
Systems	Execution Context Reuse	w/o Dual-VM Architecture	w/o Identical Kernel	Trust to Guest Kernel	High Code Coverage	Fully Automated	Memory Introspection	Disk Introspection	Guest Management	Process Monitoring
VIRTUOSO	X	✓	X	✓	X	X	✓	X	X	X
VMST	✓	X	X	✓	✓	✓	✓	X	X	X
EXTERIOR	✓	X	X	✓	✓	✓	✓	X	✓	X
PROCESSIMPLANTING	✓	✓	✓	X	✓	✓	X	X	✓	X
PROCESSOUTGRAFTING	✓	X	✓	✓	✓	✓	X	X	X	✓
GEARS	✓	✓	✓	X	✓	X	✓	✓	✓	✓
HYPERSHELL	✓	✓	✓	X	✓	✓	✓	✓	✓	✓

Virtualization (Hypervisor) Layer Applications

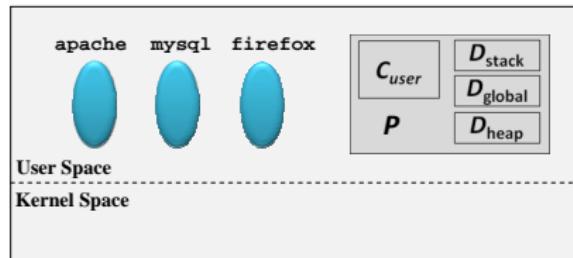


- Virtual Machine Introspection
- Virtual Machine (Re)Configuration, Repair
- Automated Out-of-VM Management

Two Approaches to bridging the semantic gap

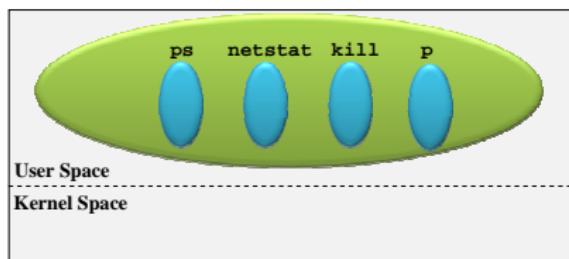


Secure VM (SVM)

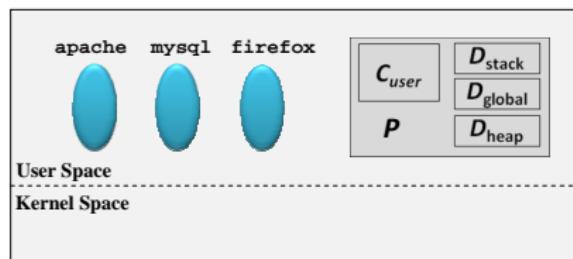


Guest VM (GVM)

Two Approaches to bridging the semantic gap



Secure VM (SVM)

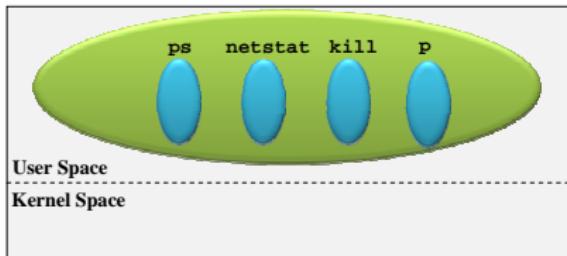


Guest VM (GVM)

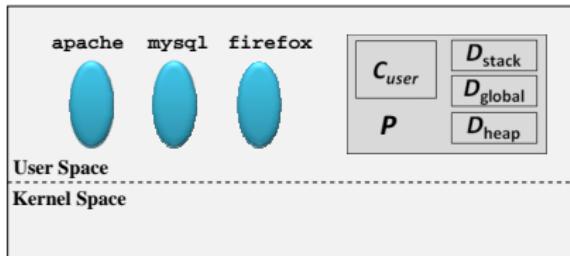
Reusing (**legacy**) **binary code** with a trusted sibling VM to introspect the running guest VM.

- ① Redirect kernel data [SP'12, VEE'13, NDSS'14] → Fine-grained, slower performance
- ② Redirect system call execution [USENIX ATC'14] → More practical, fast performance

For Cloud Developers: No Gap, Everything is Native



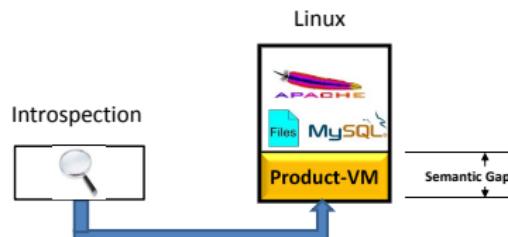
Secure VM (SVM)



Guest VM (GVM)

In-VM getpid Program

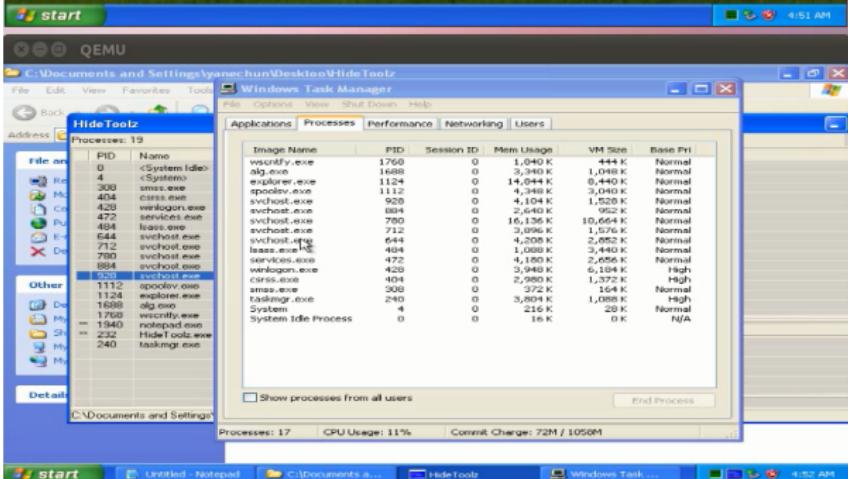
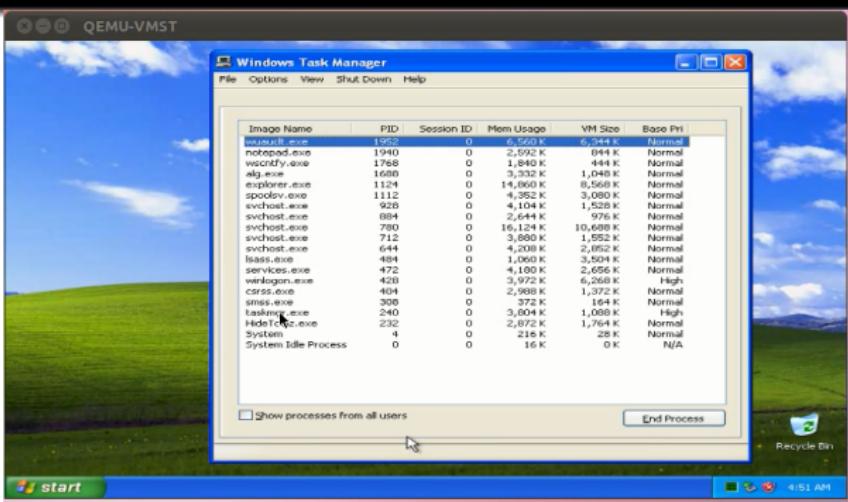
```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main()
5 {
6     printf("pid=%d\n",getpid());
7     return 0;
8 }
```



References

- 1 VM Space Traveling: Automatically Bridging the Semantic Gap in Virtual Machine Introspection via Online Kernel Data Redirection (*IEEE Symposium on Security and Privacy [SP'12]*)
- 2 EXTERIOR: Using A Dual-VM Enabled External Shell for Guest-OS Introspection, Configuration, and Recovery (*ACM SIGPLAN/SIGOPS International Conference on Virtual Execution Environments [VEE'13]*)
- 3 Hybrid-Bridge: Efficiently Bridging the Semantic-Gap in Virtual Machine Introspection via Decoupled Execution and Training Memoization (*Network and Distributed System Symposium [NDSS'14]*)
- 4 HyperShell: A Practical Hypervisor Layer Guest OS Shell for Automated In-VM Management (*USENIX Annual Technical Conference [ATC'14]*)

<http://www.utdallas.edu/~zhiqiang.lin/s3.html>



Background
ooooo

Related Work
oooooooooo

Binary Code Reuse
ooooooo

Evaluation
oooooooooooo

Summary
oooooo●

